

# Big Data Visual Analytics (CS 661)

Instructor: Soumya Dutta

Department of Computer Science and Engineering

Indian Institute of Technology Kanpur (IITK)

email: [soumyad@cse.iitk.ac.in](mailto:soumyad@cse.iitk.ac.in)

# Acknowledgements

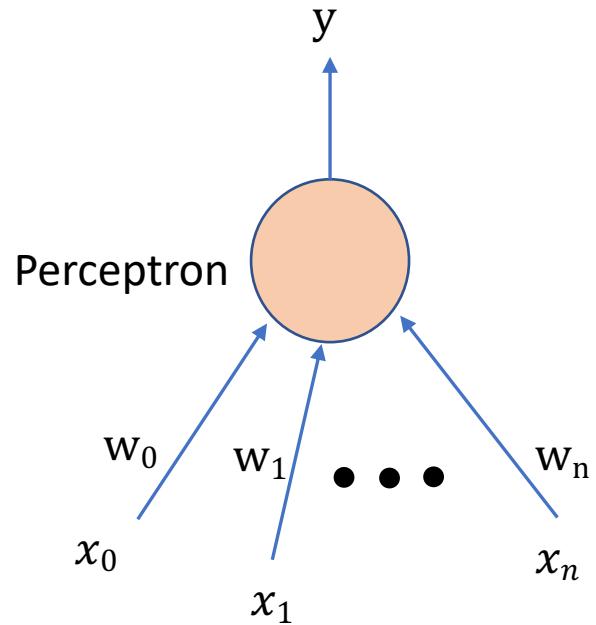
- Prof. Mitesh Khapra (IIT Madras)

# Study Materials for Lecture 22

- Dive into Deep Learning (<https://d2l.ai/index.html>)
- Deep Learning by Prof. Mitesh Khapra, IITM, NPTEL Course
  - [https://www.youtube.com/playlist?list=PLyqSpQzTE6M9gCgajvQbc68Hk\\_JKGBAYT](https://www.youtube.com/playlist?list=PLyqSpQzTE6M9gCgajvQbc68Hk_JKGBAYT)
- NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis
- On the Spectral Bias of Neural Networks, Rahaman et al.
- Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains, Tancik et al.

# Basics of Feed Forward Neural Network

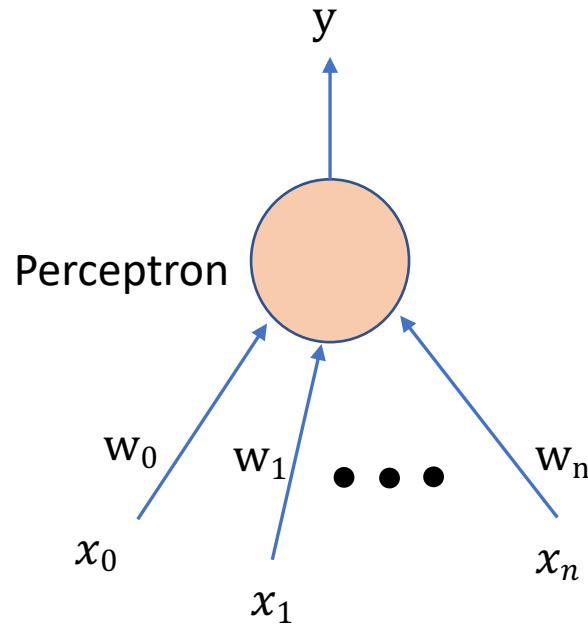
# Perceptron and Sigmoid Neuron



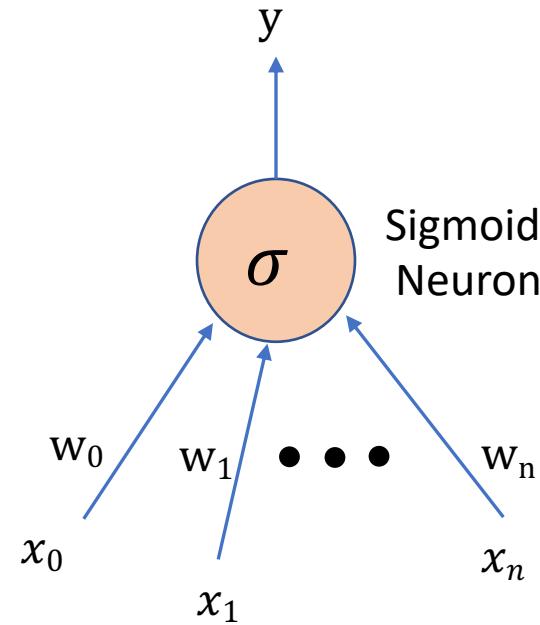
$$y = 1 \text{ if } \sum_{i=0}^n w_i * x_i \geq 0$$

$$= 0 \text{ if } \sum_{i=0}^n w_i * x_i < 0$$

# Perceptron and Sigmoid (Logistic) Neuron

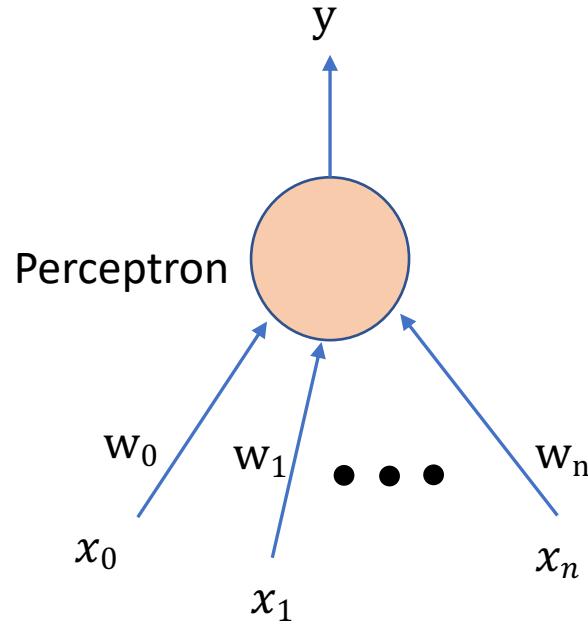


$$y = 1 \text{ if } \sum_{i=0}^n w_i * x_i \geq 0$$
$$= 0 \text{ if } \sum_{i=0}^n w_i * x_i < 0$$



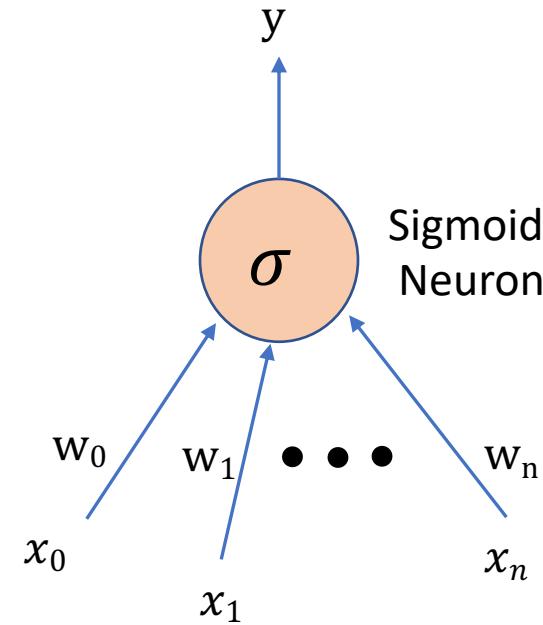
$$y = \frac{1}{1 + e^{-(\sum_{i=0}^n w_i * x_i)}}$$

# Perceptron and Sigmoid (Logistic) Neuron

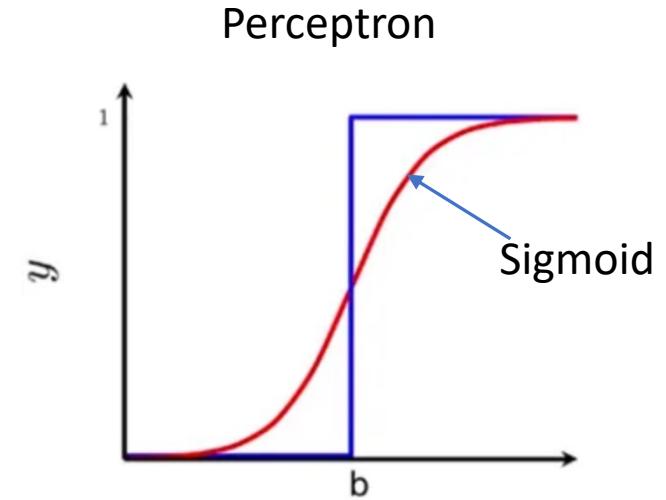


$$y = 1 \text{ if } \sum_{i=0}^n w_i * x_i \geq 0$$

$$= 0 \text{ if } \sum_{i=0}^n w_i * x_i < 0$$



$$y = \frac{1}{1 + e^{-(\sum_{i=0}^n w_i * x_i)}}$$

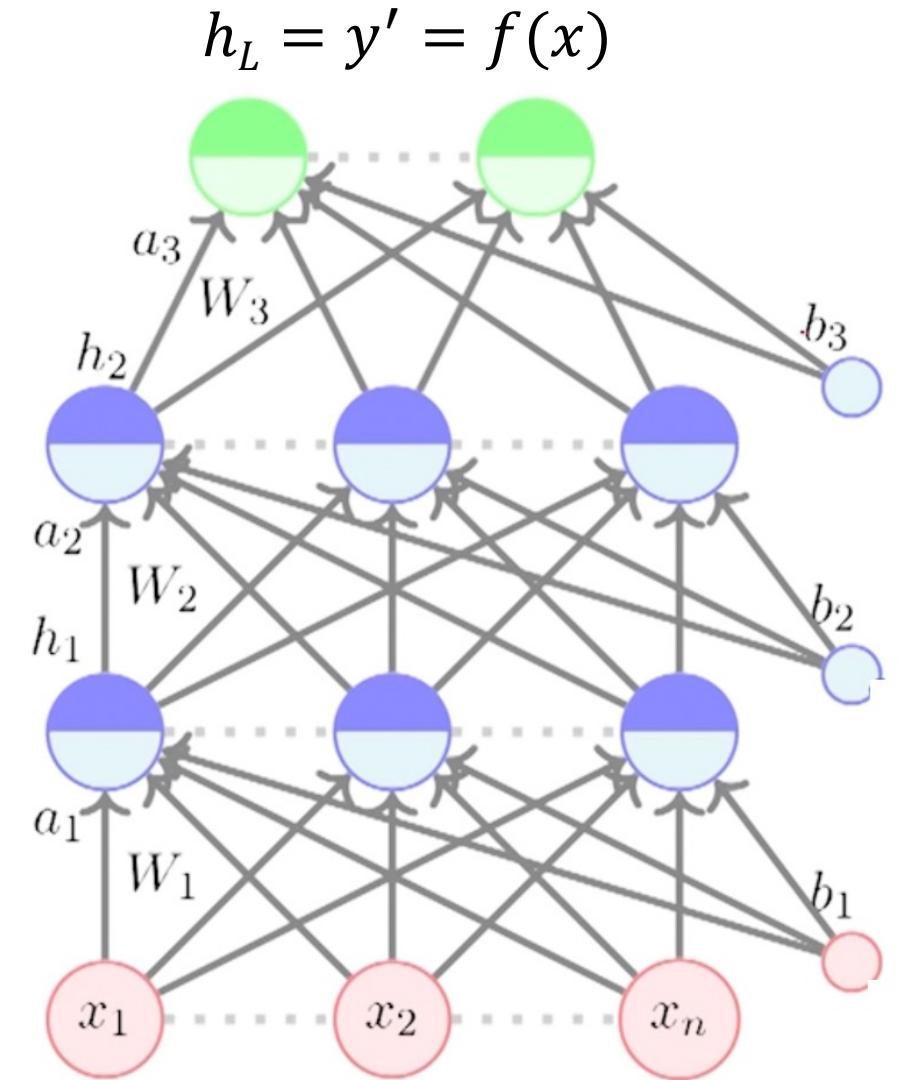


# Feed Forward Neural Network

- The input to the network is an  $n$ -dimensional vector
- The network contains  $L - 1$  hidden layers, having  $n$  neurons in each hidden layer
- There is one output layer containing  $k$  neurons
- Each neuron in hidden layer can be thought of having two parts:
  - pre-activation and activation
- The input layer can be called 0-th layer, and the output layer can be called  $L$ -th layer

# Feed Forward Neural Network

- This network has 2 hidden layers
- ‘ $a$ ’ and ‘ $h$ ’ are vectors of length  $n$
- Dimension of  $h_L$  is  $k$  (as  $k$  neurons at output layer)
- Dimension of  $W_1$  and  $W_2$  are  $n \times n$  and for  $W_3$ , dimension is  $n \times k$
- Dimension of  $b$  (bias) vectors are  $n$ , for output layer, dimension of  $b$  is  $k$



# Feed Forward Neural Network

- The pre-activation at layer  $i$  is given by

$$a_i(x) = b_i + w_i h_{i-1}(x)$$

- The activation at layer  $i$  is given by

$$h_i(x) = g(a_i(x))$$

Here  $g(\cdot)$  can be sigmoid function, which is called the activation function

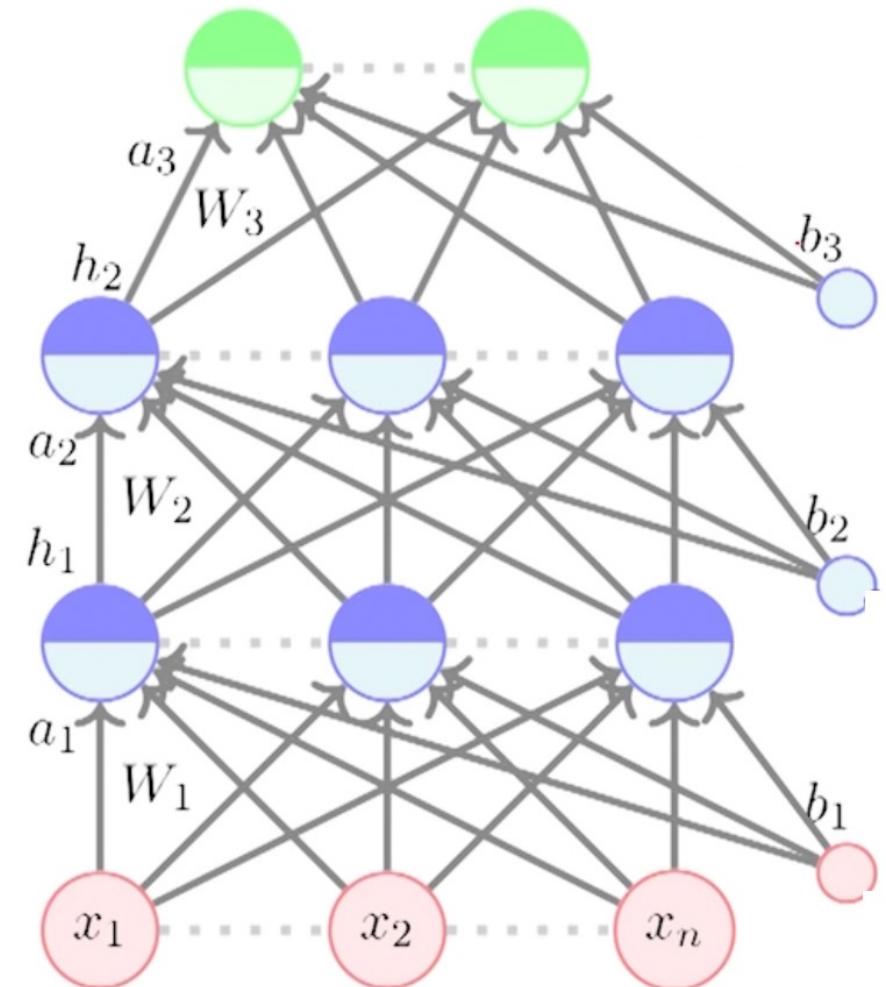
- The output at layer  $L$  is given by

$$f(x) = h_L = O(aL)$$

$O$  is the output activation function such as softmax, linear, etc.

$$y' = f(x) = O(W_3g(W_2g(W_1x + b_1) + b_2 + b_3))$$

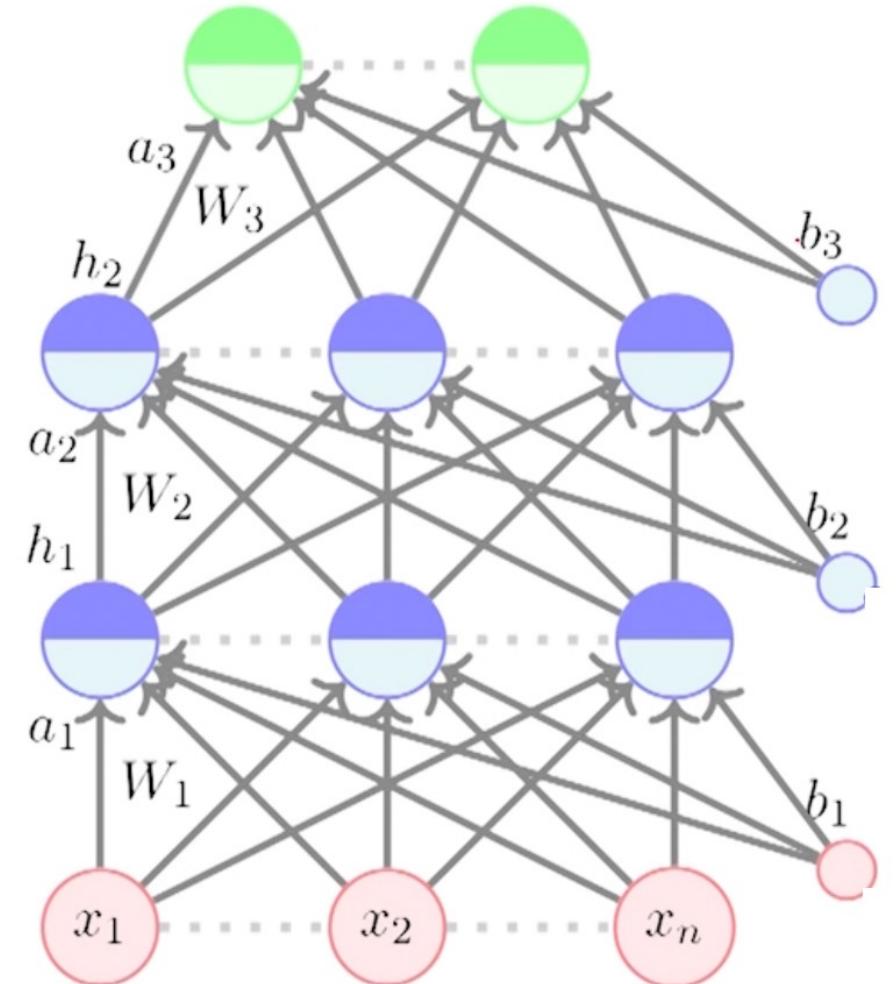
$$h_L = y' = f(x)$$



# Feed Forward Neural Network

- Data:  $\{x_i, y_i\} \forall i \in n$
- Model:  
$$y' = f(x) = O(W_3g(W_2g(W_1x + b_1) + b_2) + b_3)$$
- Parameters:  $W$  and  $b$
- Learning Algorithm: Gradient Descent with backpropagation
- Objective/Loss function: Mean squared error (MSE), Binary cross entropy, etc.

$$h_L = y' = f(x)$$



# Feed Forward Neural Network

- Gradient descent algorithm:

---

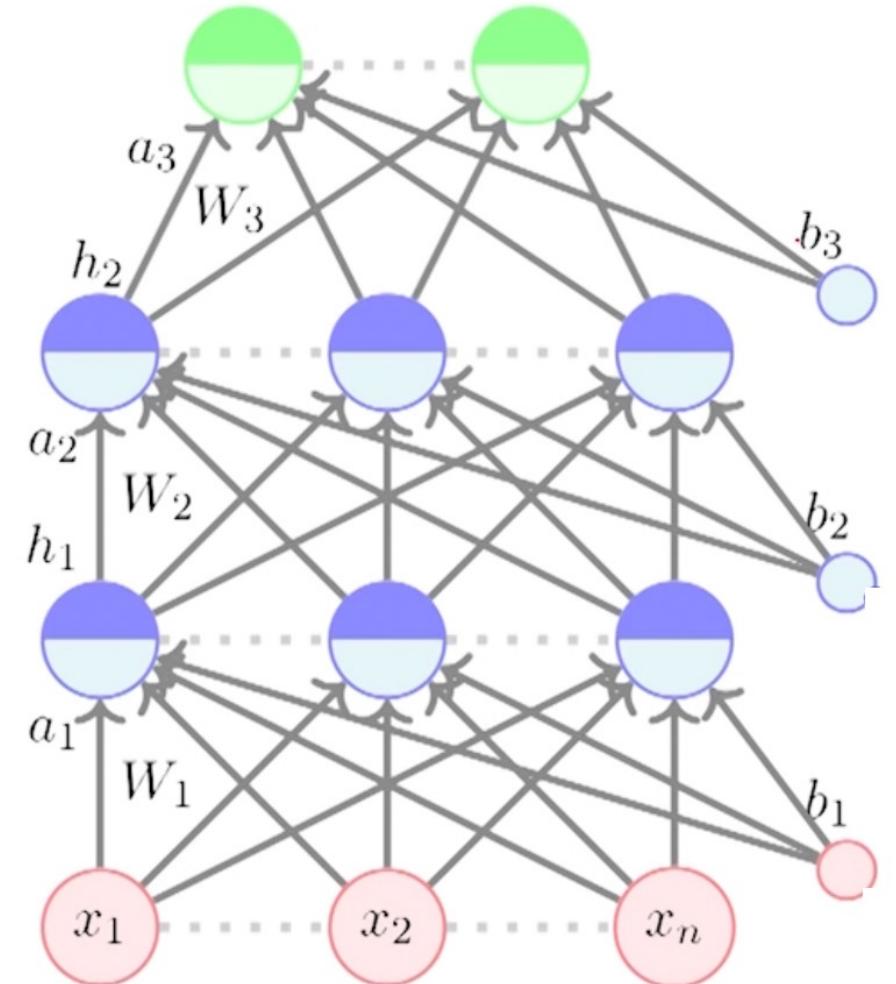
**Algorithm:** gradient\_descent()

---

```
t ← 0;  
max_iterations ← 1000;  
Initialize  $\theta_0 = [w_0, b_0]$ ;  
while  $t++ < max\_iterations$  do  
     $\theta_{t+1} \leftarrow \theta_t - \eta \nabla \theta_t$ ;  
end
```

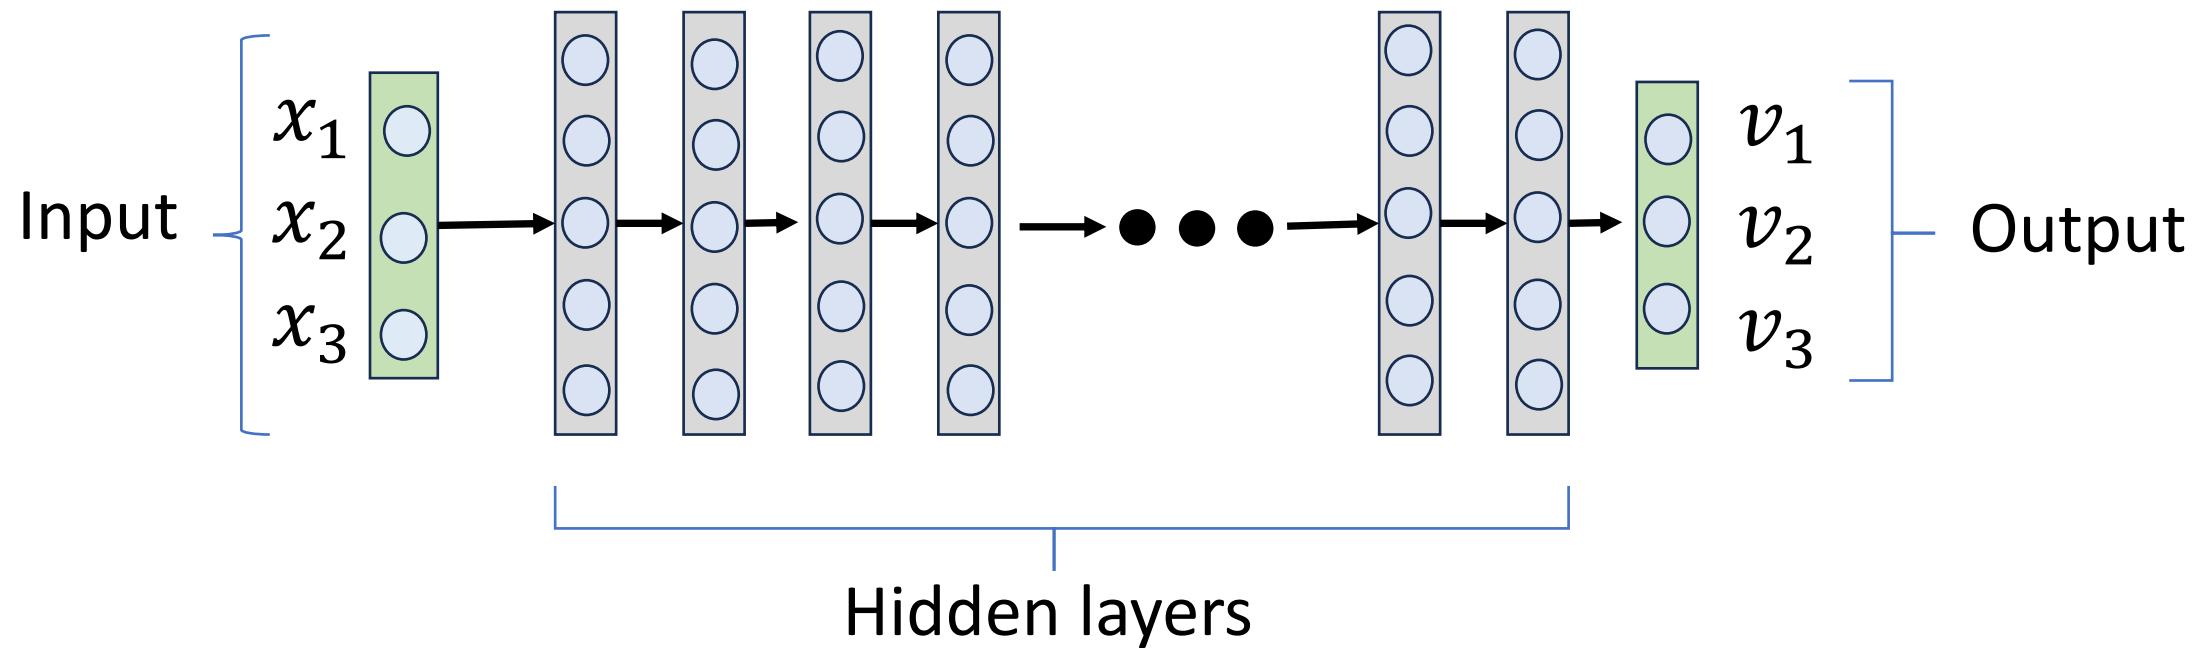
---

$$h_L = y' = f(x)$$

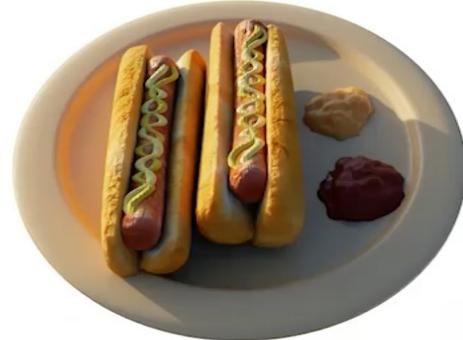


# A Multilayer Perceptron (MLP)

- A type of feed forward neural network



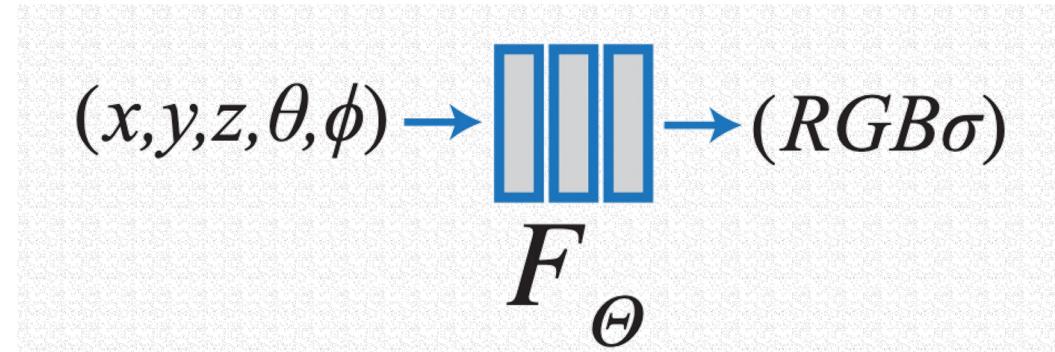
# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis



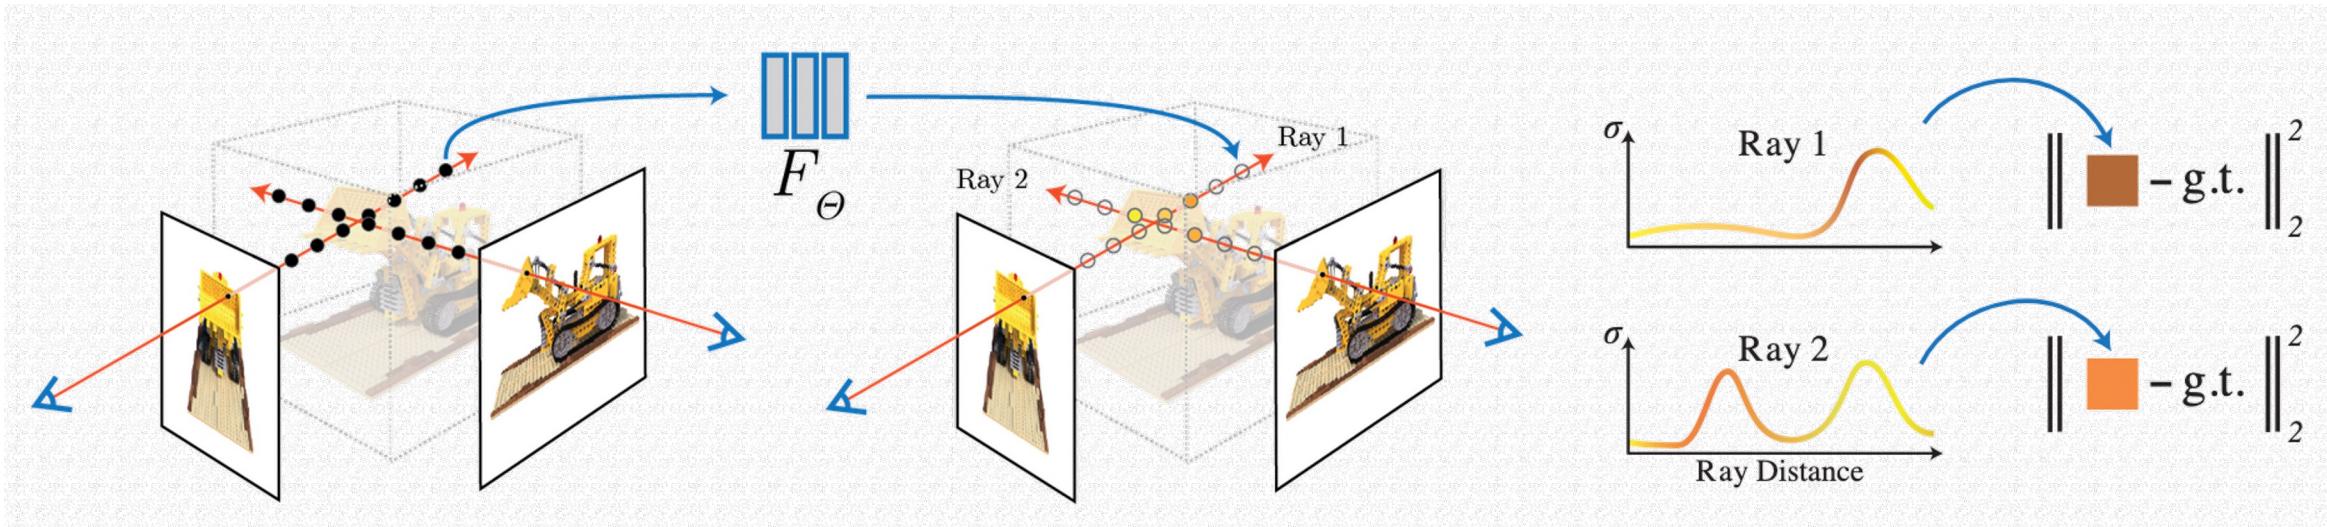
# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis



# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis



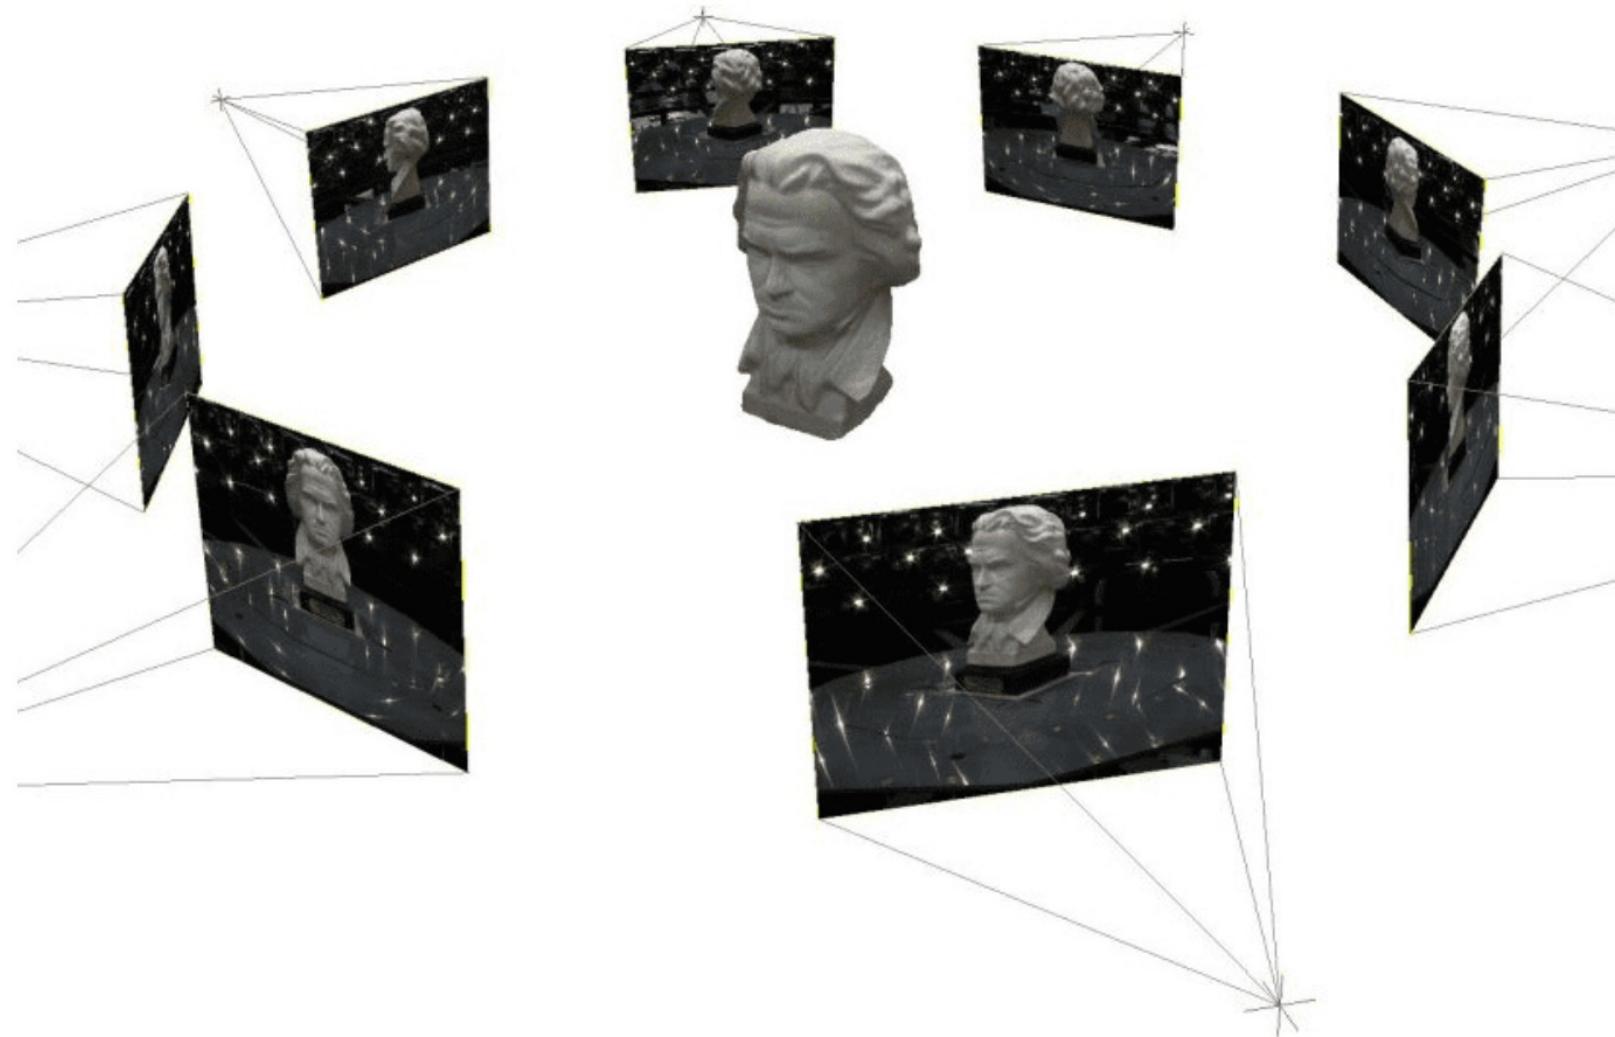
- Input: camera position and view direction
- Output: color and volume density



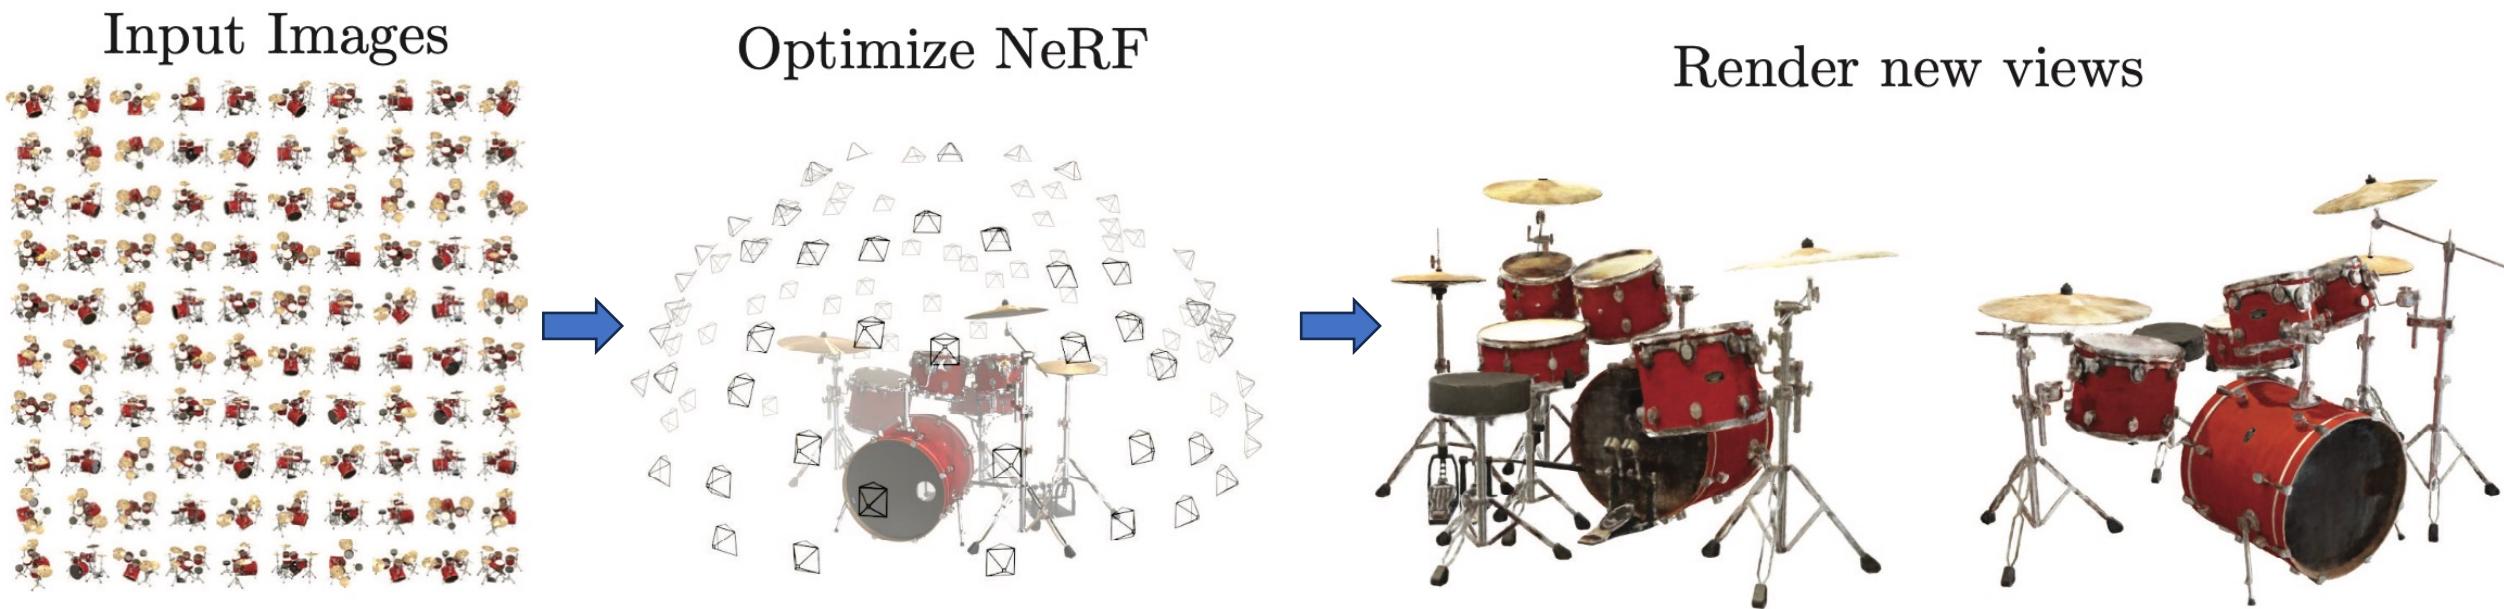
# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

- NeRF learns the radiance field in 3D volume space
  - Given a dataset containing images of a scene, their corresponding camera poses, and intrinsic parameters
  - Output: Predict the color and volume density for every viewing location and direction
- Input to the model is images of the scene from several different view directions
- After training, model can predict images for novel views with high accuracy

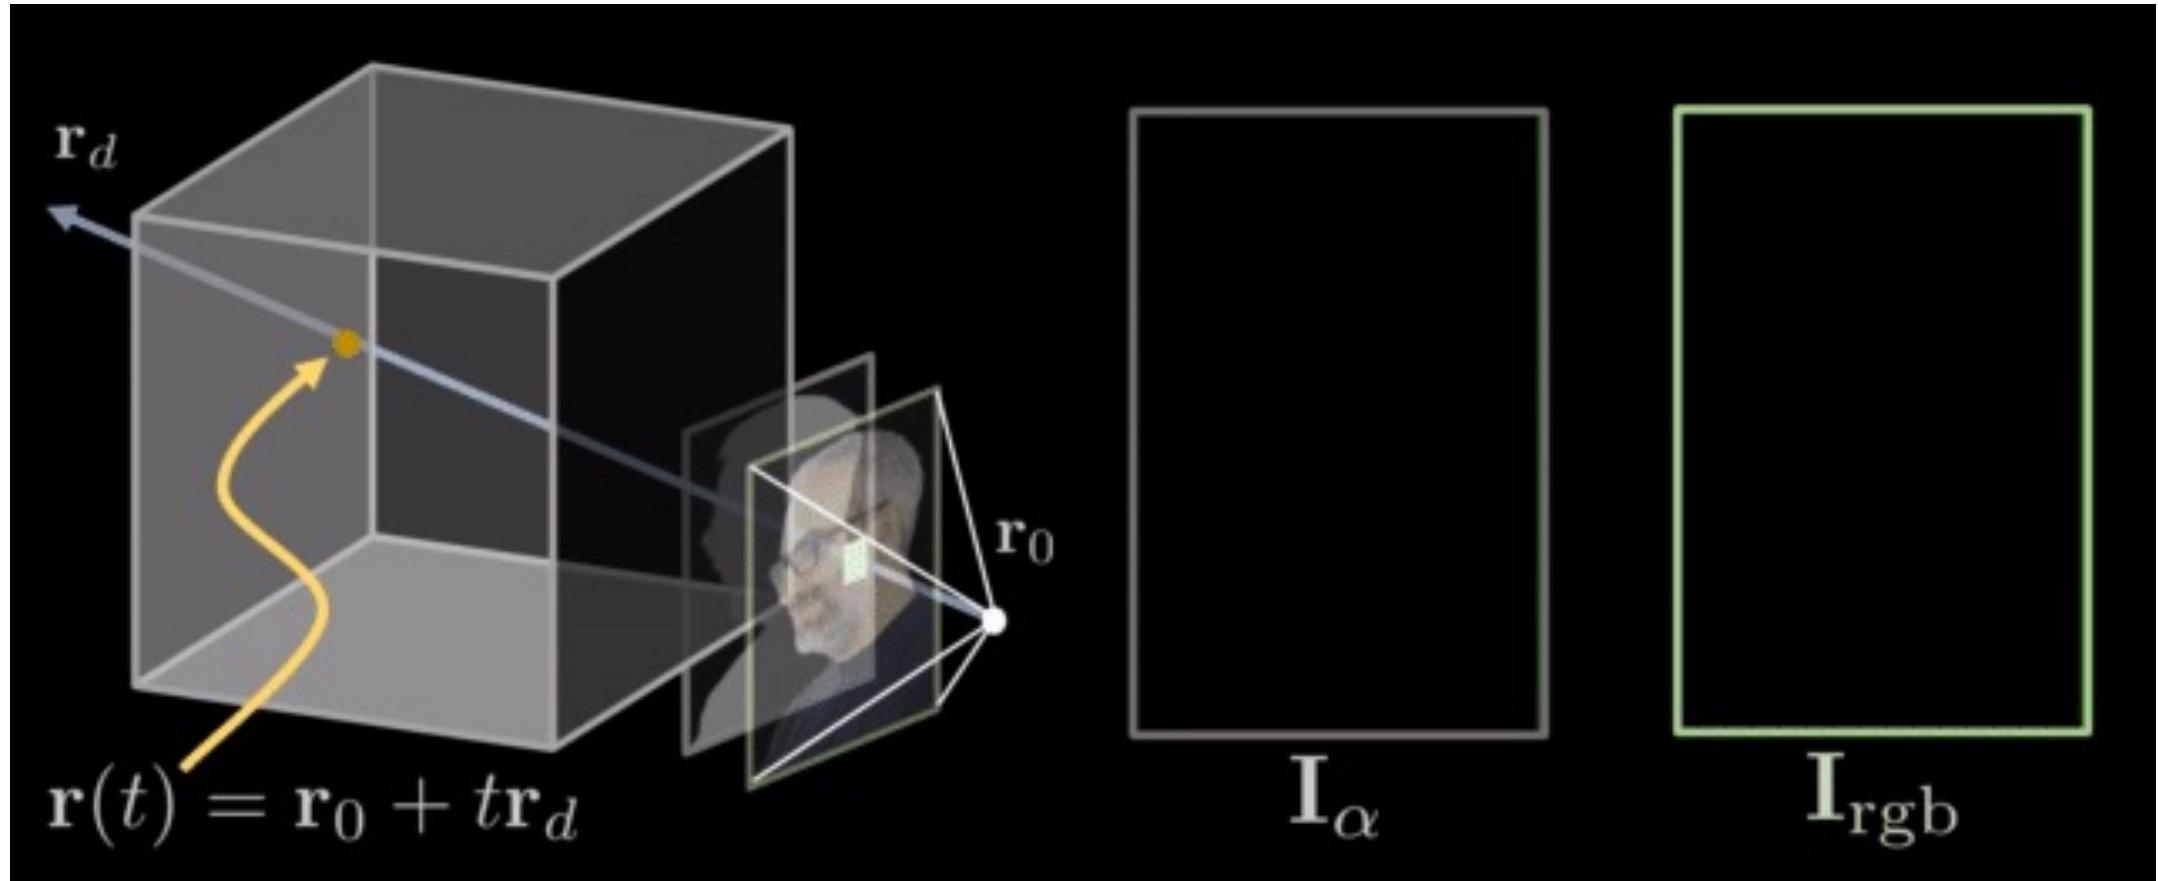
# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis



# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis



# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis



# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

- Use Volume Rendering Algorithm (Ray Casting)
- Generating a view from NeRF requires rendering all rays that pass through each pixel of the desired virtual camera

$$C(r) = \sum_{i=1}^N T_i(1 - \exp(-\sigma_i \delta_i))c_i \quad \text{where} \quad T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

Expected color of a camera ray

The relative contribution of this segment

Predicted color

Probability that nothing has blocked the ray up to this point

# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

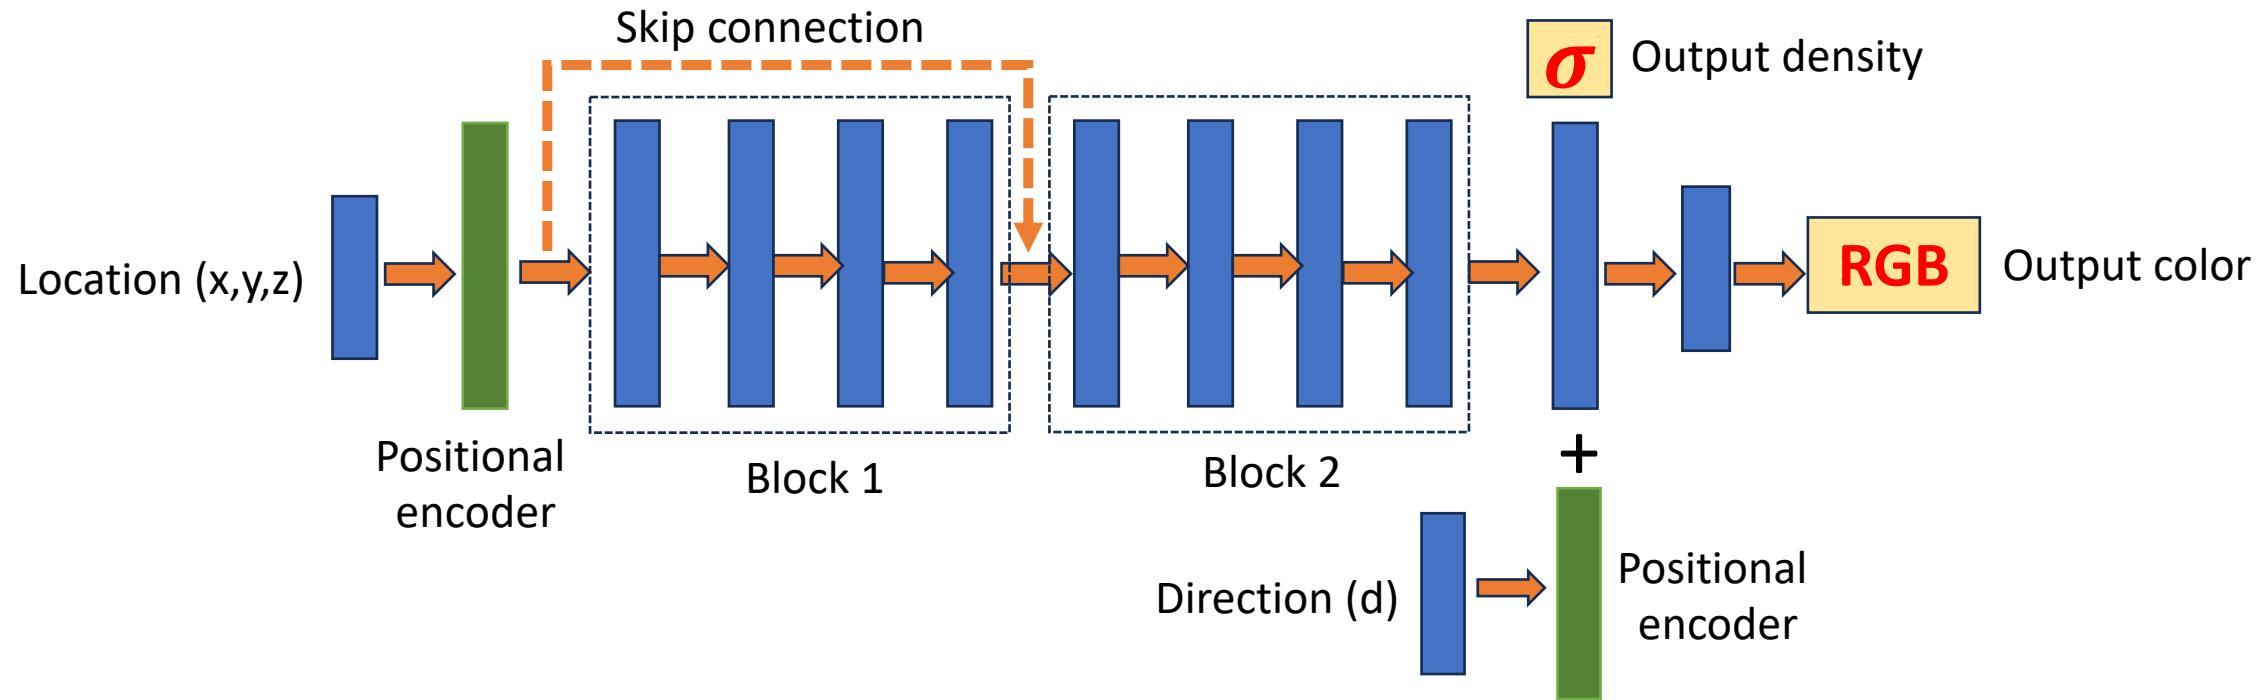
- Use Volume Rendering Algorithm (Ray Casting)
- Generating a view from NeRF requires rendering all rays that pass through each pixel of the desired virtual camera

$$C(r) = \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i \quad \text{where} \quad T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

NeRF model learns to predict

# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

- Neural Network architecture is built using simple Multi Layer Perceptron (MLP)



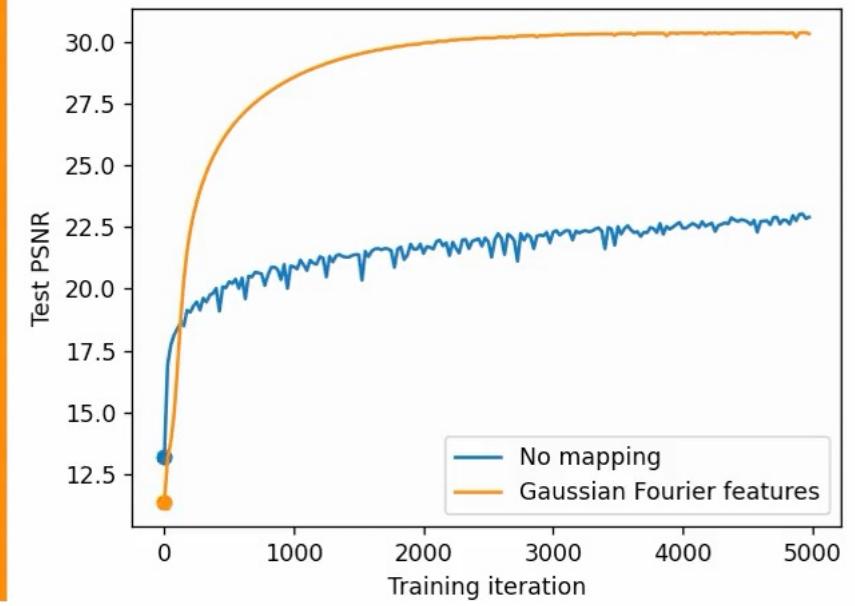
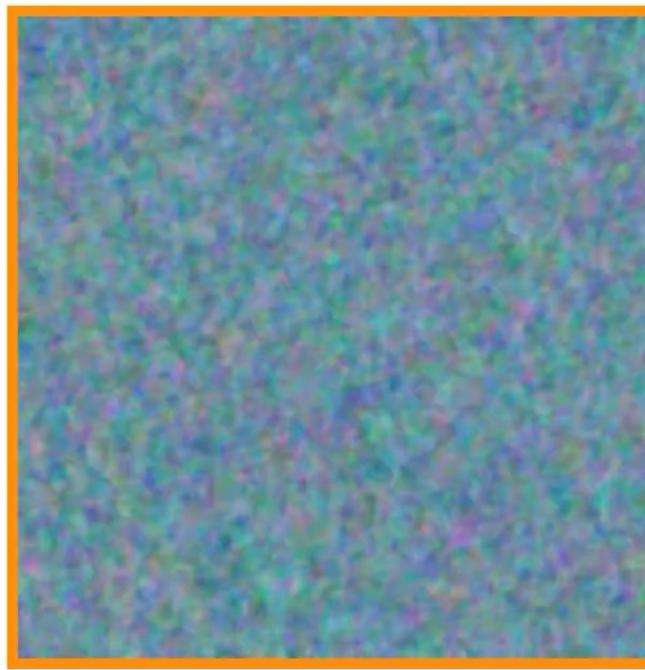
# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

- Positional Encoding:
  - Fully-connected deep networks are biased to learn low frequencies faster
  - Positional encoding helps in learning high frequency details of the data more accurately
  - Transformer architectures use similar encoding for a different purpose

$$\gamma(p) = (\sin(2^0\pi p), \cos(2^0\pi p), \dots, \sin(2^{L-1}\pi p), \cos(2^{L-1}\pi p))$$

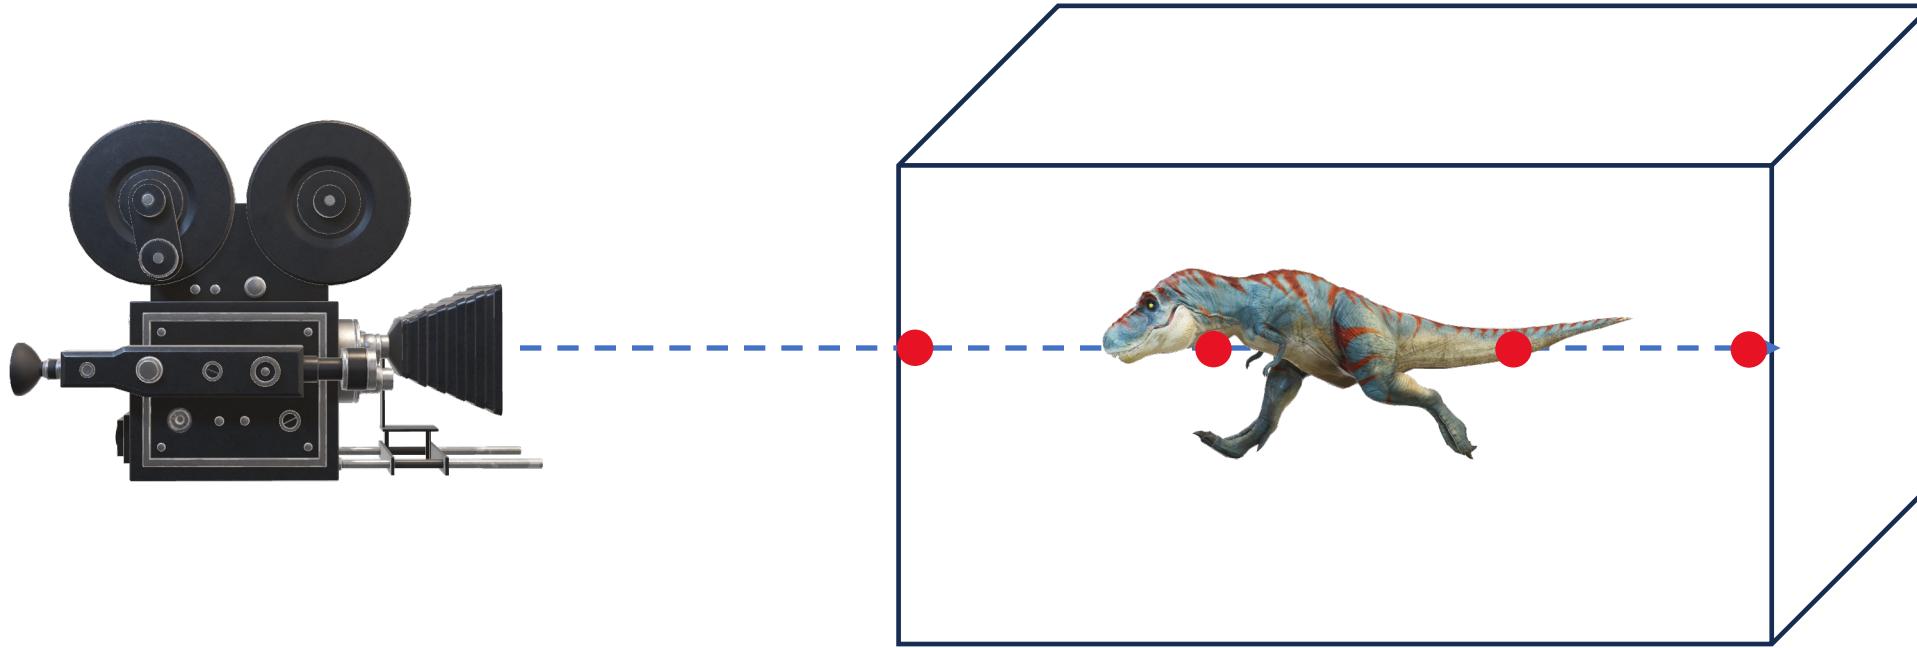
# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

- Positional Encoding:
  - Fully-connected deep networks are biased to learn low frequencies faster
  - Positional encoding helps in learning high frequency details of the data more accurately



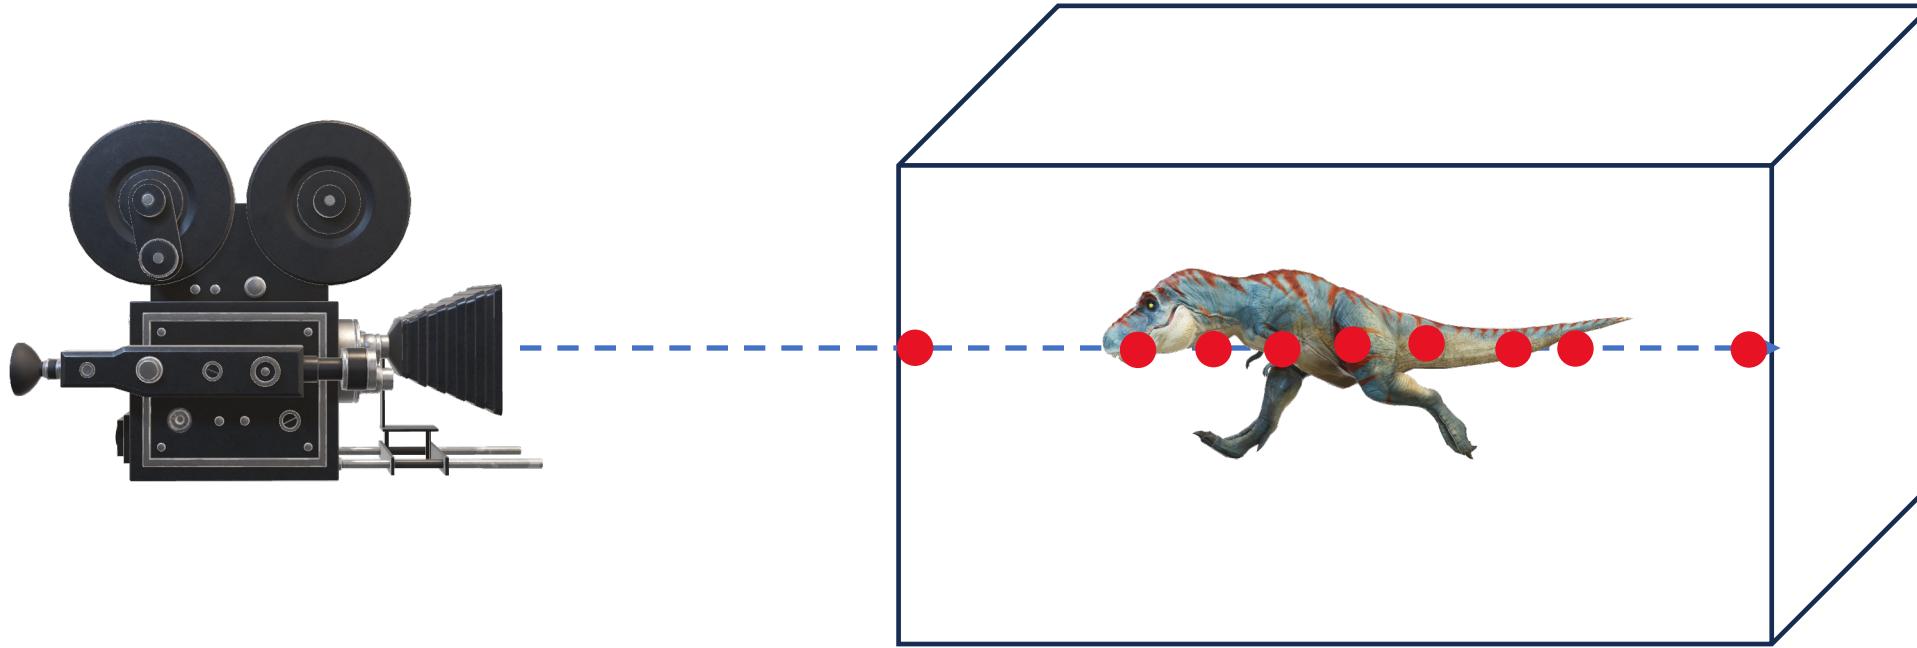
# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

- Hierarchical Volume Sampling:
  - Inefficient to integrate over empty
  - Allocate samples proportionally to their expected effect on the final rendering

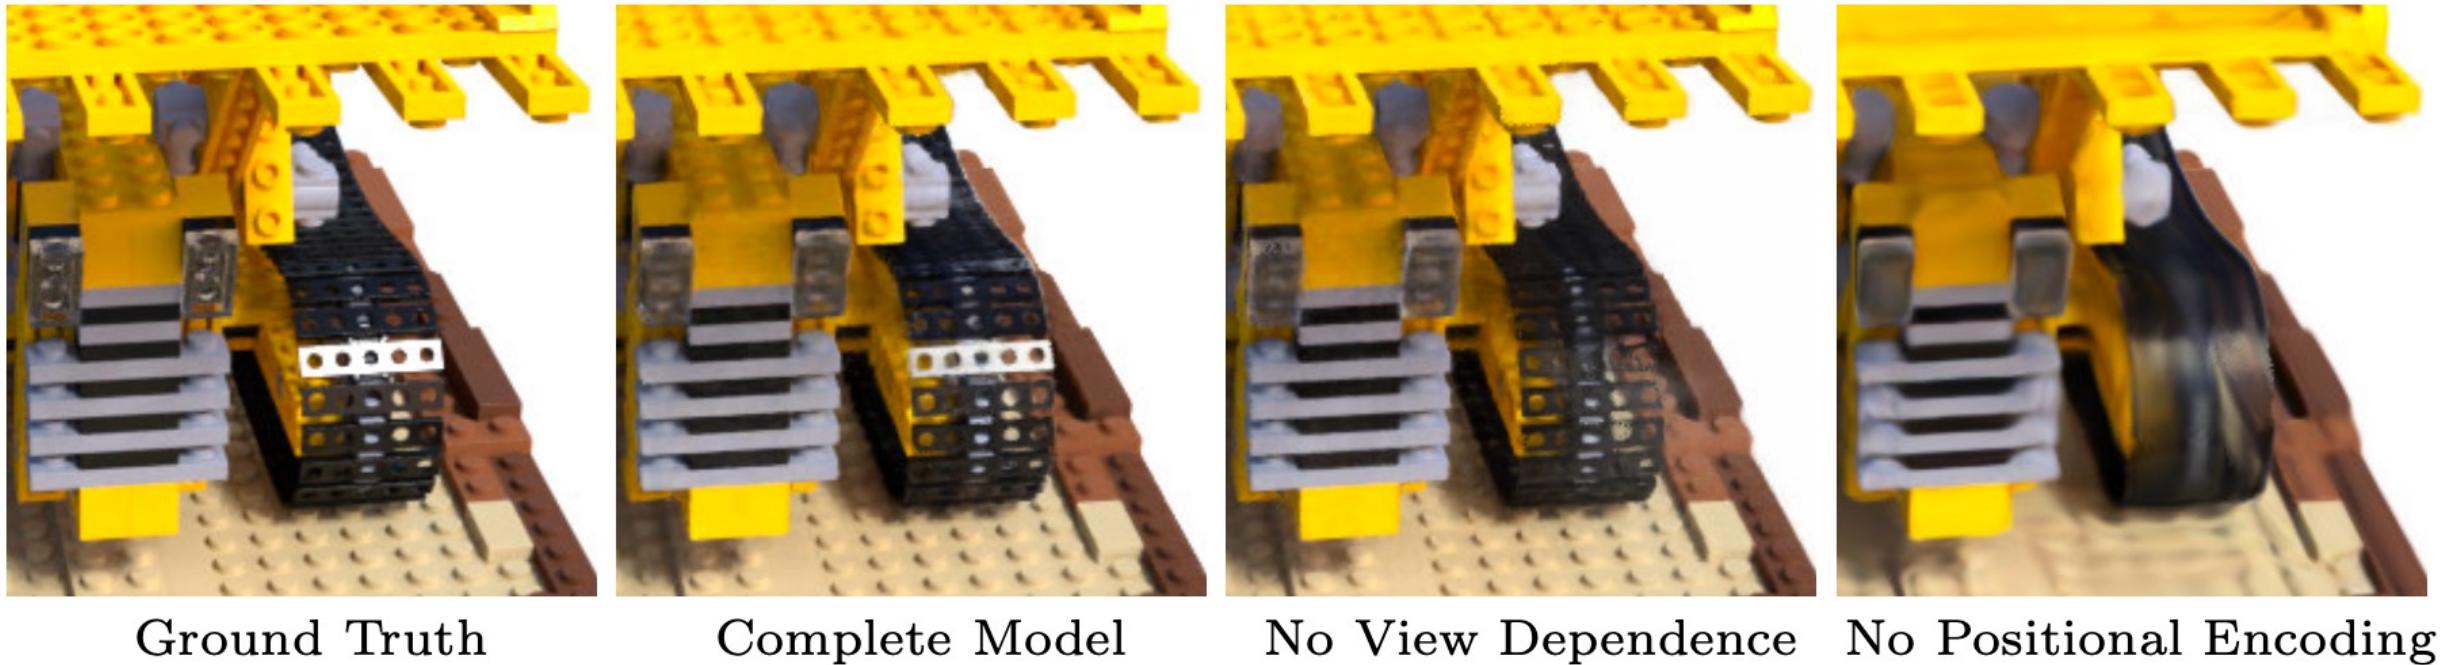


# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

- Hierarchical Volume Sampling:
  - Inefficient to integrate over empty
  - Allocate samples proportionally to their expected effect on the final rendering



# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis



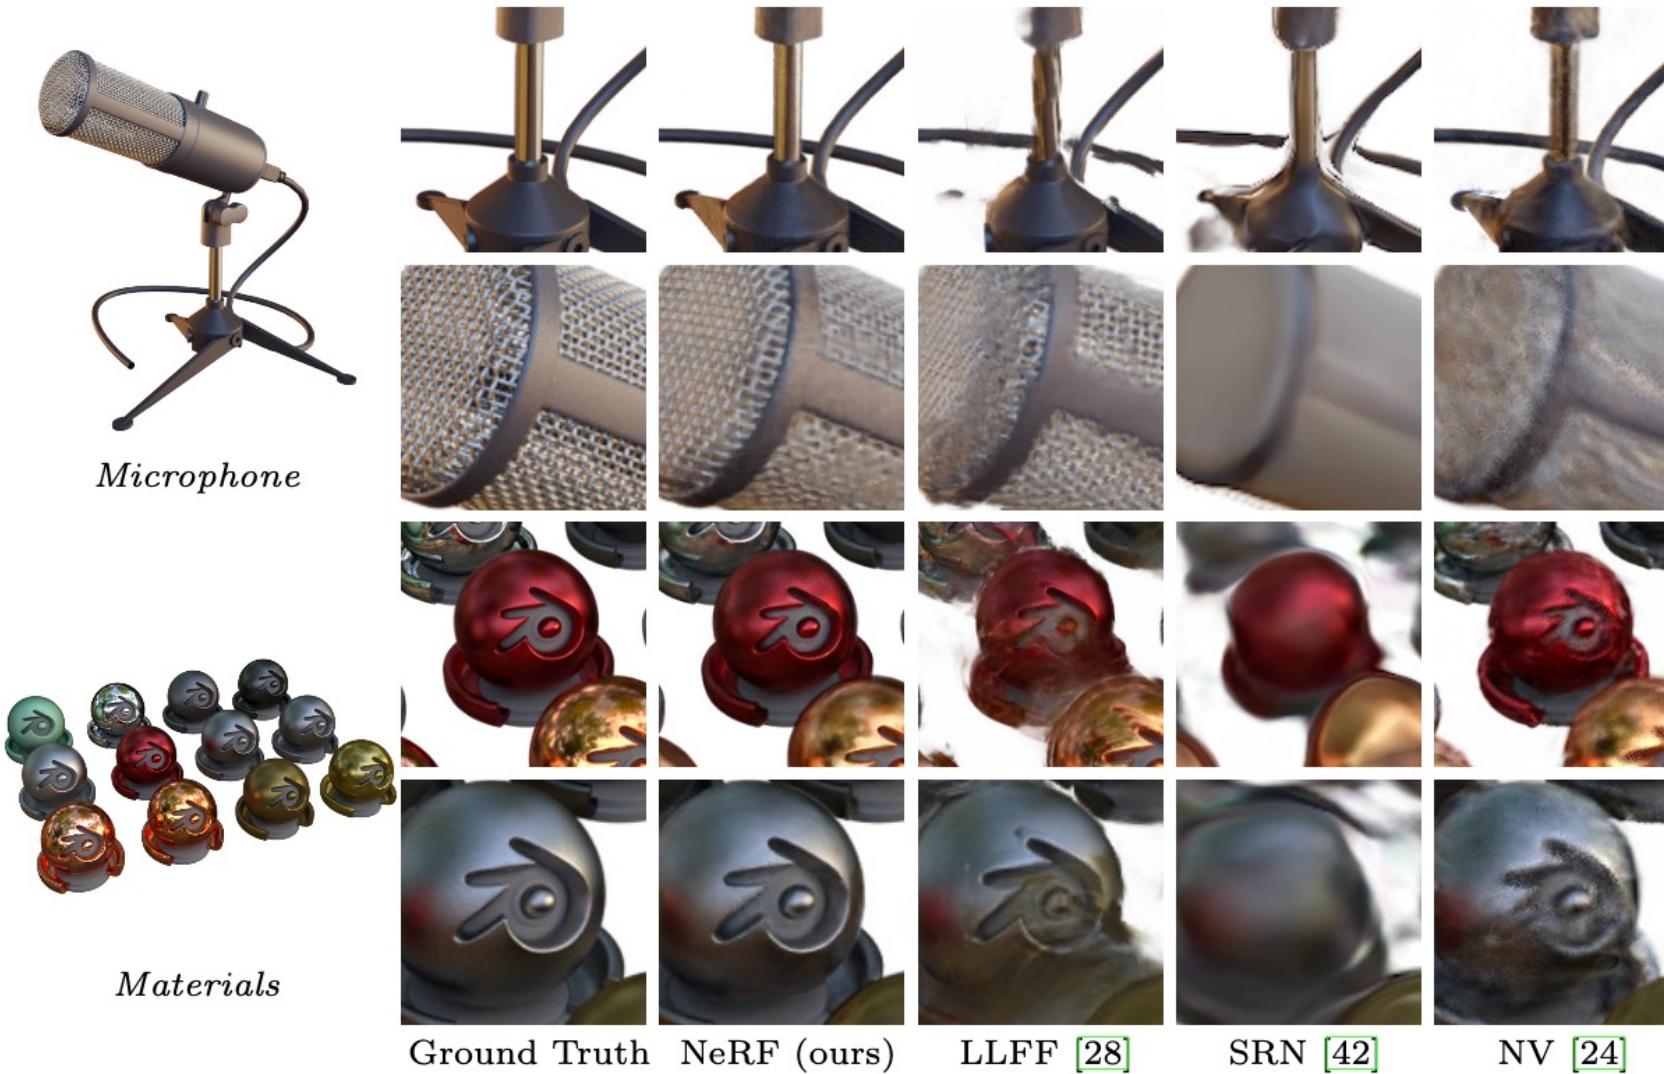
Ground Truth

Complete Model

No View Dependence

No Positional Encoding

# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis



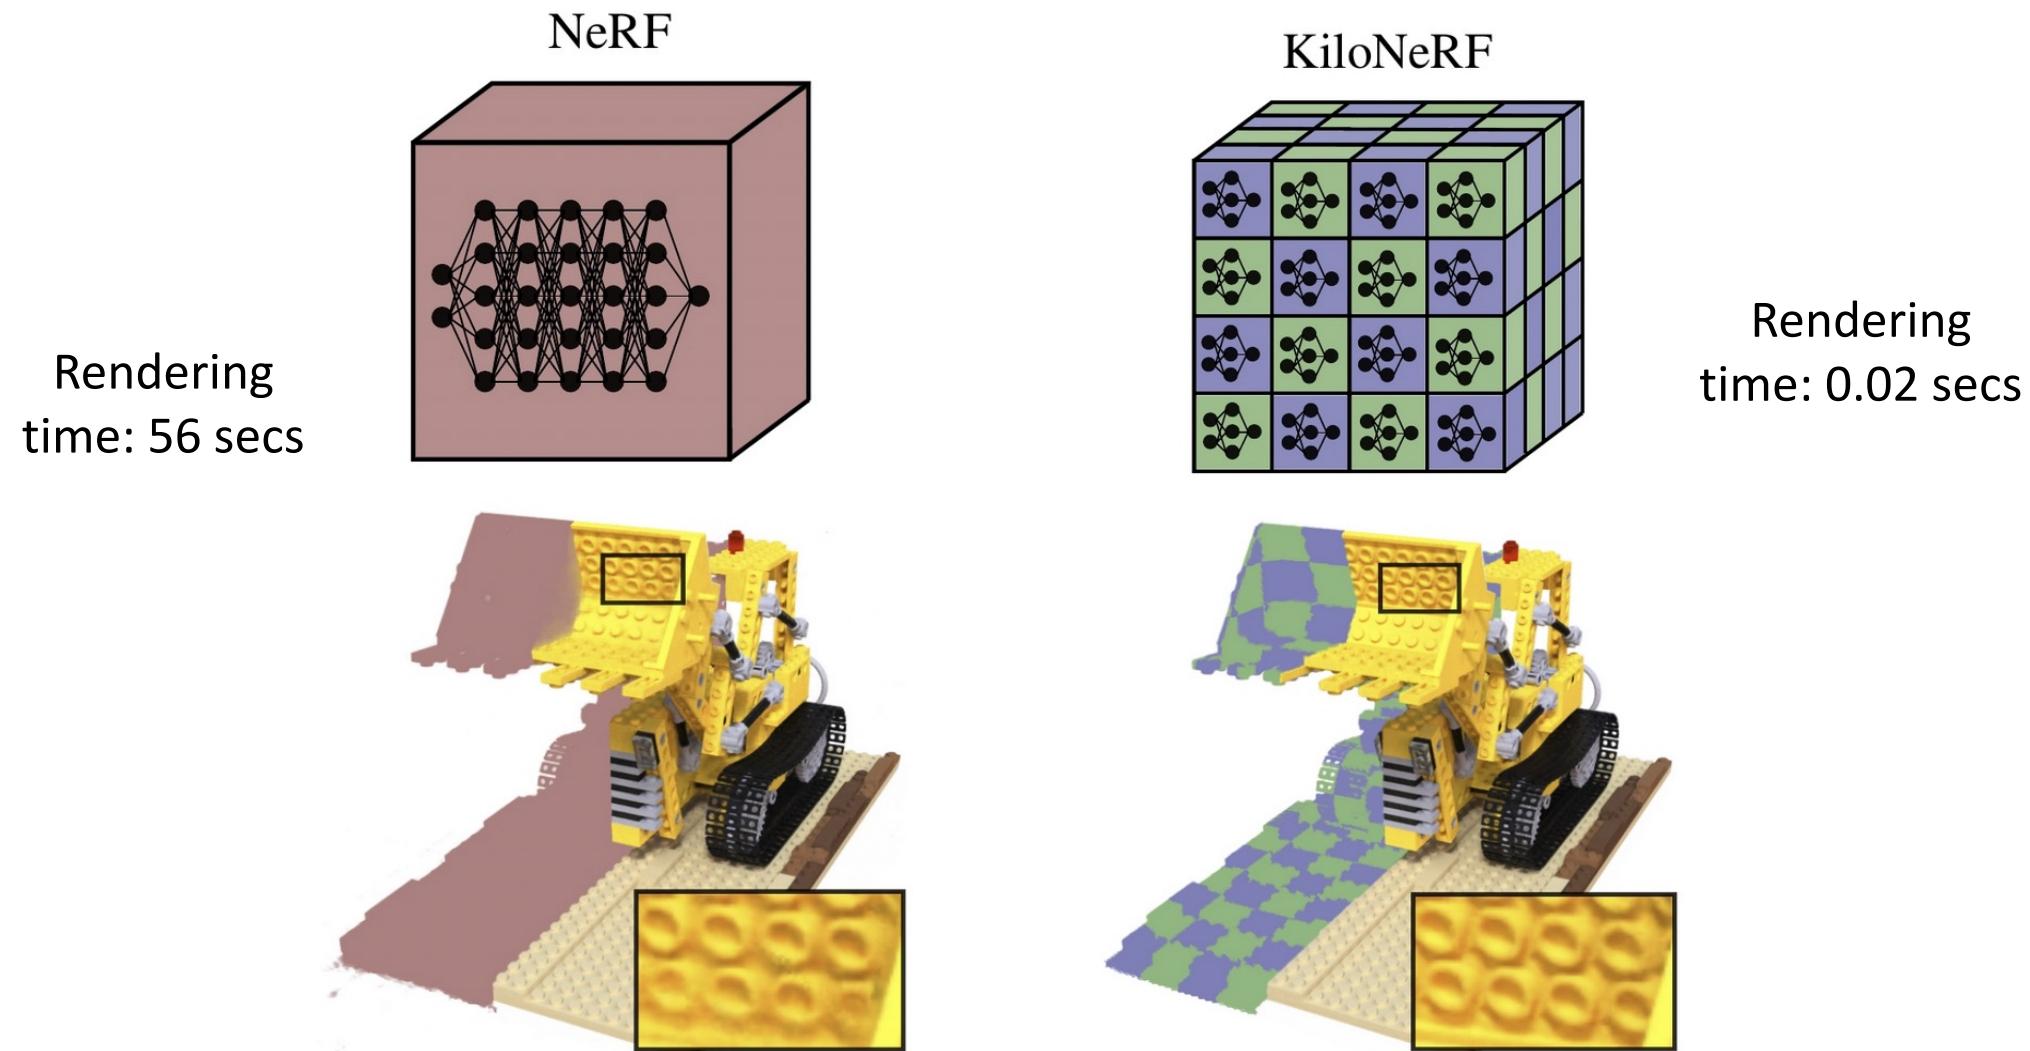
# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis



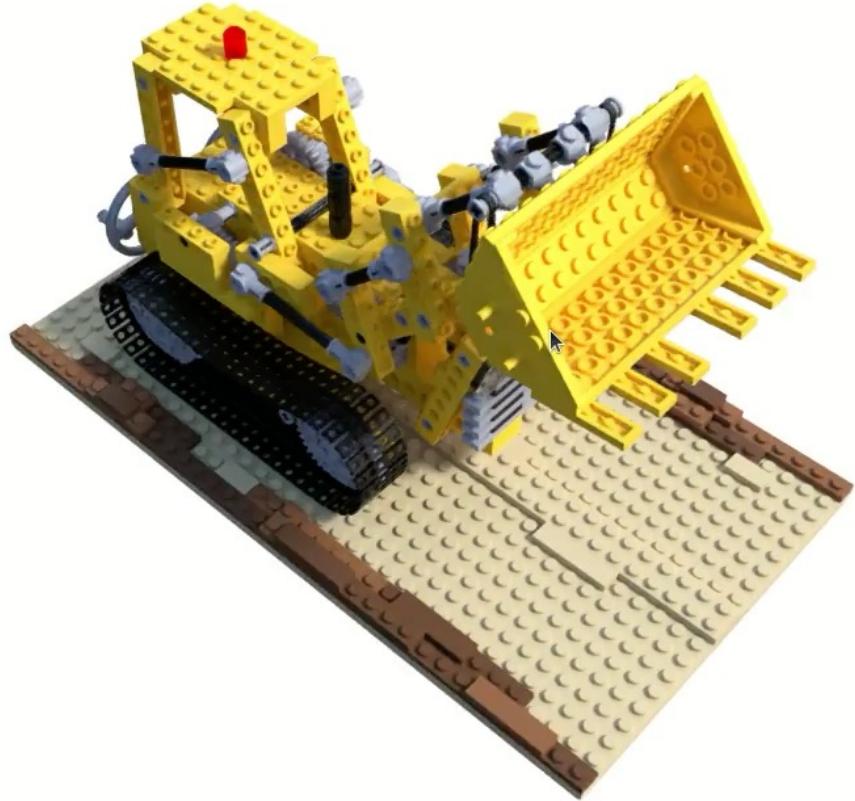
# NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis

	Input	#Im.	$L$	( $N_c$ , $N_f$ )	PSNR↑	SSIM↑	LPIPS↓
1) No PE, VD, H	$xyz$	100	-	(256, -)	26.67	0.906	0.136
2) No Pos. Encoding	$xyz\theta\phi$	100	-	(64, 128)	28.77	0.924	0.108
3) No View Dependence	$xyz$	100	10	(64, 128)	27.66	0.925	0.117
4) No Hierarchical	$xyz\theta\phi$	100	10	(256, -)	30.06	0.938	0.109
5) Far Fewer Images	$xyz\theta\phi$	25	10	(64, 128)	27.78	0.925	0.107
6) Fewer Images	$xyz\theta\phi$	50	10	(64, 128)	29.79	0.940	0.096
7) Fewer Frequencies	$xyz\theta\phi$	100	5	(64, 128)	30.59	0.944	0.088
8) More Frequencies	$xyz\theta\phi$	100	15	(64, 128)	30.81	0.946	0.096
9) Complete Model	$xyz\theta\phi$	100	10	(64, 128)	<b>31.01</b>	<b>0.947</b>	<b>0.081</b>

# KiloNeRF: Speeding up Neural Radiance Fields with Thousands of Tiny MLPs



# KiloNeRF: Speeding up Neural Radiance Fields with Thousands of Tiny MLPs



# Other works based on Vanilla NeRF

- <https://github.com/awesome-NeRF/awesome-NeRF>

## Papers

---

- [NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis](#), Mildenhall et al., ECCV 2020 | [github](#) | [bibtex](#)
- ▼ Faster Inference
  - [Learning Neural Transmittance for Efficient Rendering of Reflectance Fields](#), Mohammad Shafiei et al., BMVC 2021 | [bibtex](#)
  - [Neural Sparse Voxel Fields](#), Liu et al., NeurIPS 2020 | [github](#) | [bibtex](#)
  - [AutoInt: Automatic Integration for Fast Neural Volume Rendering](#), Lindell et al., CVPR 2021 | [github](#) | [bibtex](#)
  - [DeRF: Decomposed Radiance Fields](#), Rebain et al. Arxiv 2020 | [bibtex](#)
  - [DONeRF: Towards Real-Time Rendering of Compact Neural Radiance Fields using Depth Oracle Networks](#), Neff et al., CGF 2021 | [github](#) | [bibtex](#)
  - [FastNeRF: High-Fidelity Neural Rendering at 200FPS](#), Garbin et al., Arxiv 2021 | [bibtex](#)
  - [KiloNeRF: Speeding up Neural Radiance Fields with Thousands of Tiny MLPs](#), Reiser et al., ICCV 2021 | [github](#) | [bibtex](#)
  - [PlenOctrees for Real-time Rendering of Neural Radiance Fields](#), Yu et al., Arxiv 2021 | [github](#) | [bibtex](#)
  - [Mixture of Volumetric Primitives for Efficient Neural Rendering](#), Lombardi et al., SIGGRAPH 2021 | [bibtex](#)
  - [Light Field Networks: Neural Scene Representations with Single-Evaluation Rendering](#), Sitzmann et al., Arxiv 2021 | [github](#) | [bibtex](#)
  - [RT-NeRF: Real-Time On-Device Neural Radiance Fields Towards Immersive AR/VR Rendering](#), Li et al., ICCAD 2022 | [bibtex](#)
  - [ENeRF: Efficient Neural Radiance Fields for Interactive Free-viewpoint Video](#), Lin et al., SIGGRAPH 2022 | [github](#) | [bibtex](#)

# Spectral Bias of (ReLU) MLPs

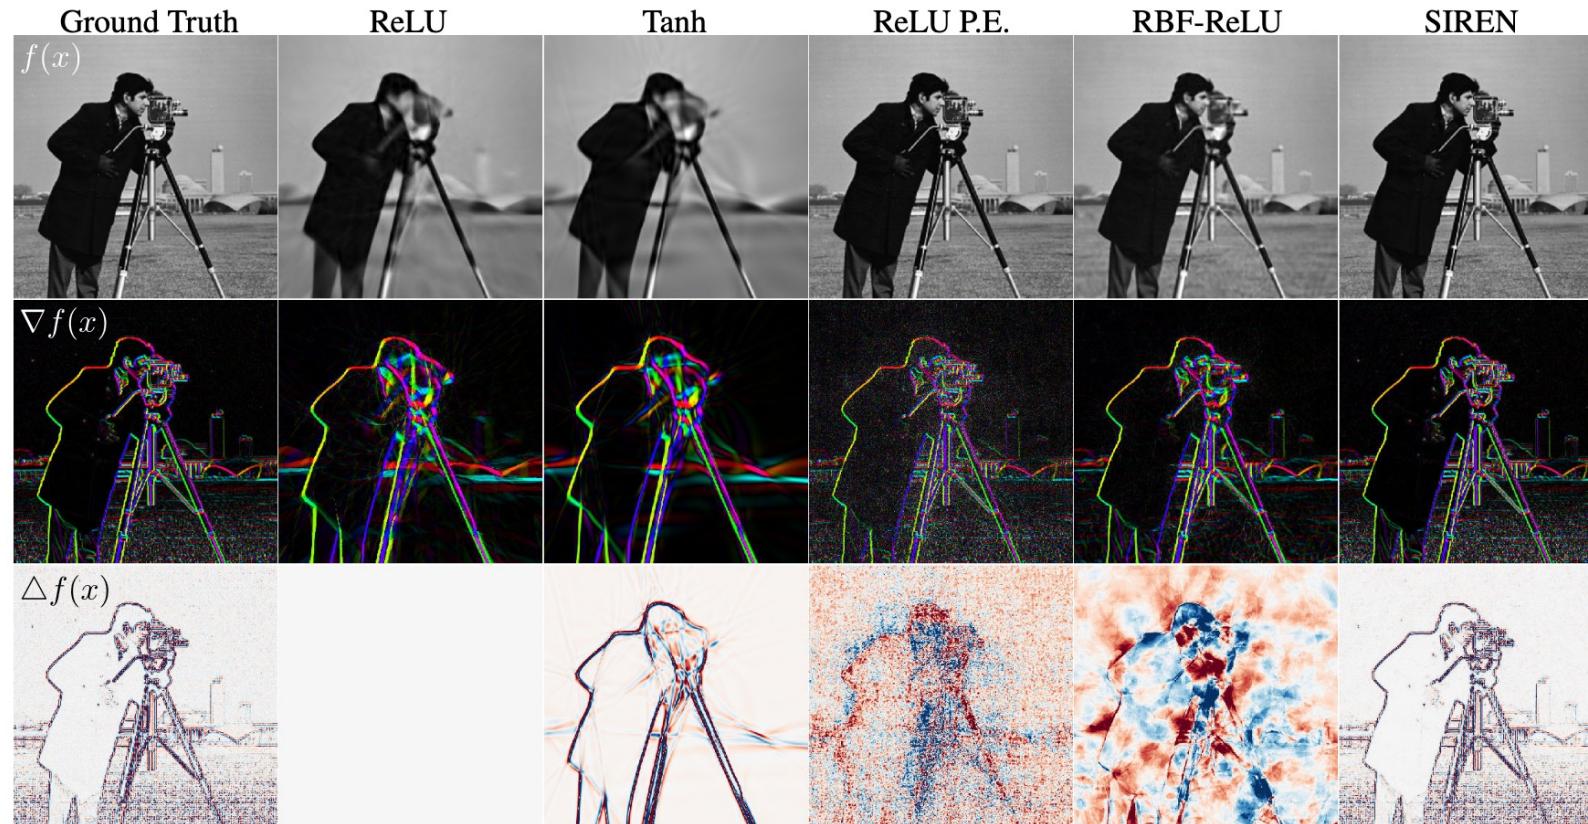
- Deep Networks prioritize learning simple patterns that usually generalize well across data samples
  - On the Spectral Bias of Neural Networks, Rahaman et al.
  - Prioritizes learning low frequency information
  - The spectral response decays quickly
  - Learning of high frequency information takes longer time and often deeper network architecture
- Solutions
  - Use Fourier feature mapping, i.e., positional encoding
    - Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains, Tancik et al.
  - Use sinusoidal activation function

# MLP with Sinusoidal Activation Function

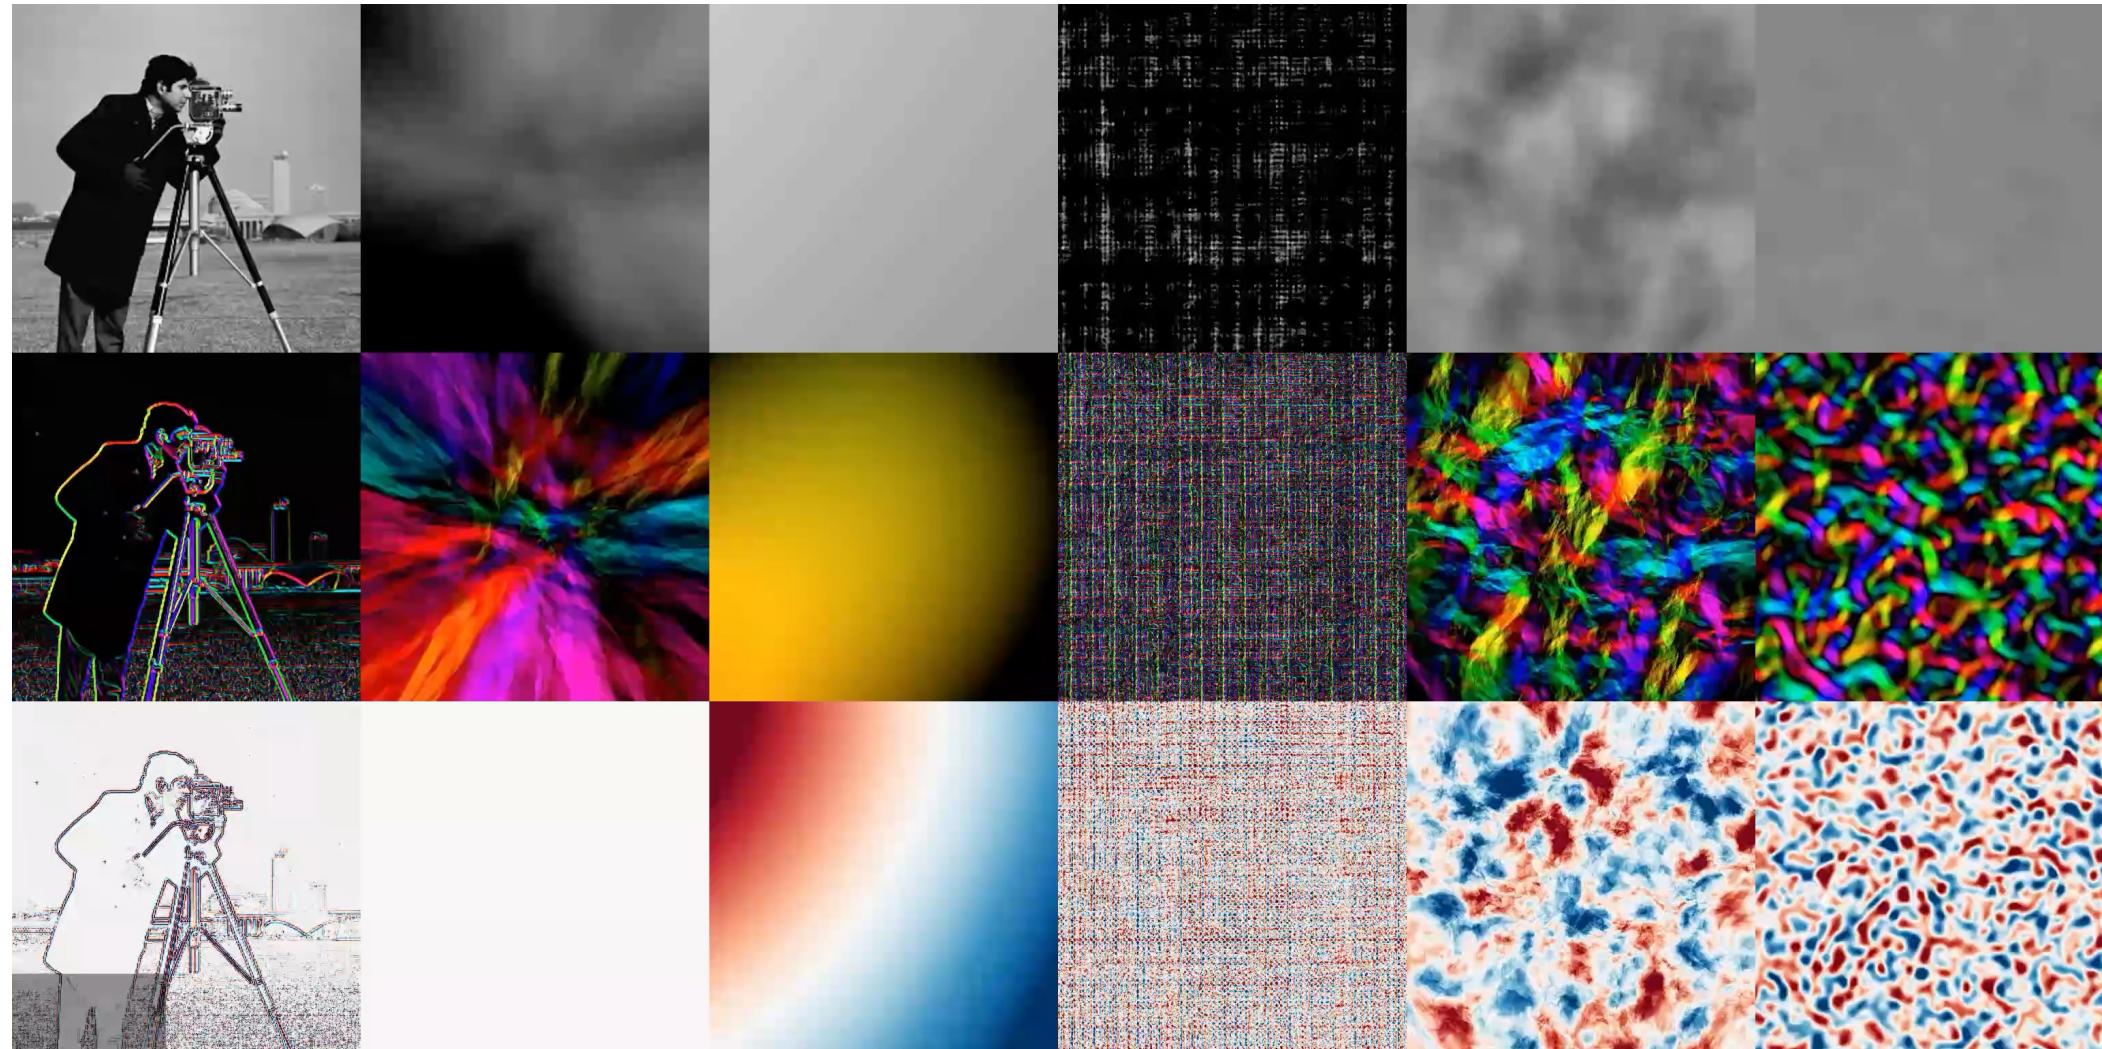
- A good fit for images, 3D field type data that can be periodic and contain information of various types of frequencies
  - Textures, edges, fine details
- Learning in a rich Fourier spectrum space
  - Model can approximate functions with high-frequency components *without* needing deep networks
- Stable Gradients
  - sine function provides stable and informative gradients
  - If initialized carefully, gradients do not vanish or explode

# SIREN: SI<sub>n</sub>usoidal REpresentation Network

- A multi-layer perceptron network with sinusoidal activation function
- Excellent to model coordinate-based data sets
- Can learn higher order gradients of the data

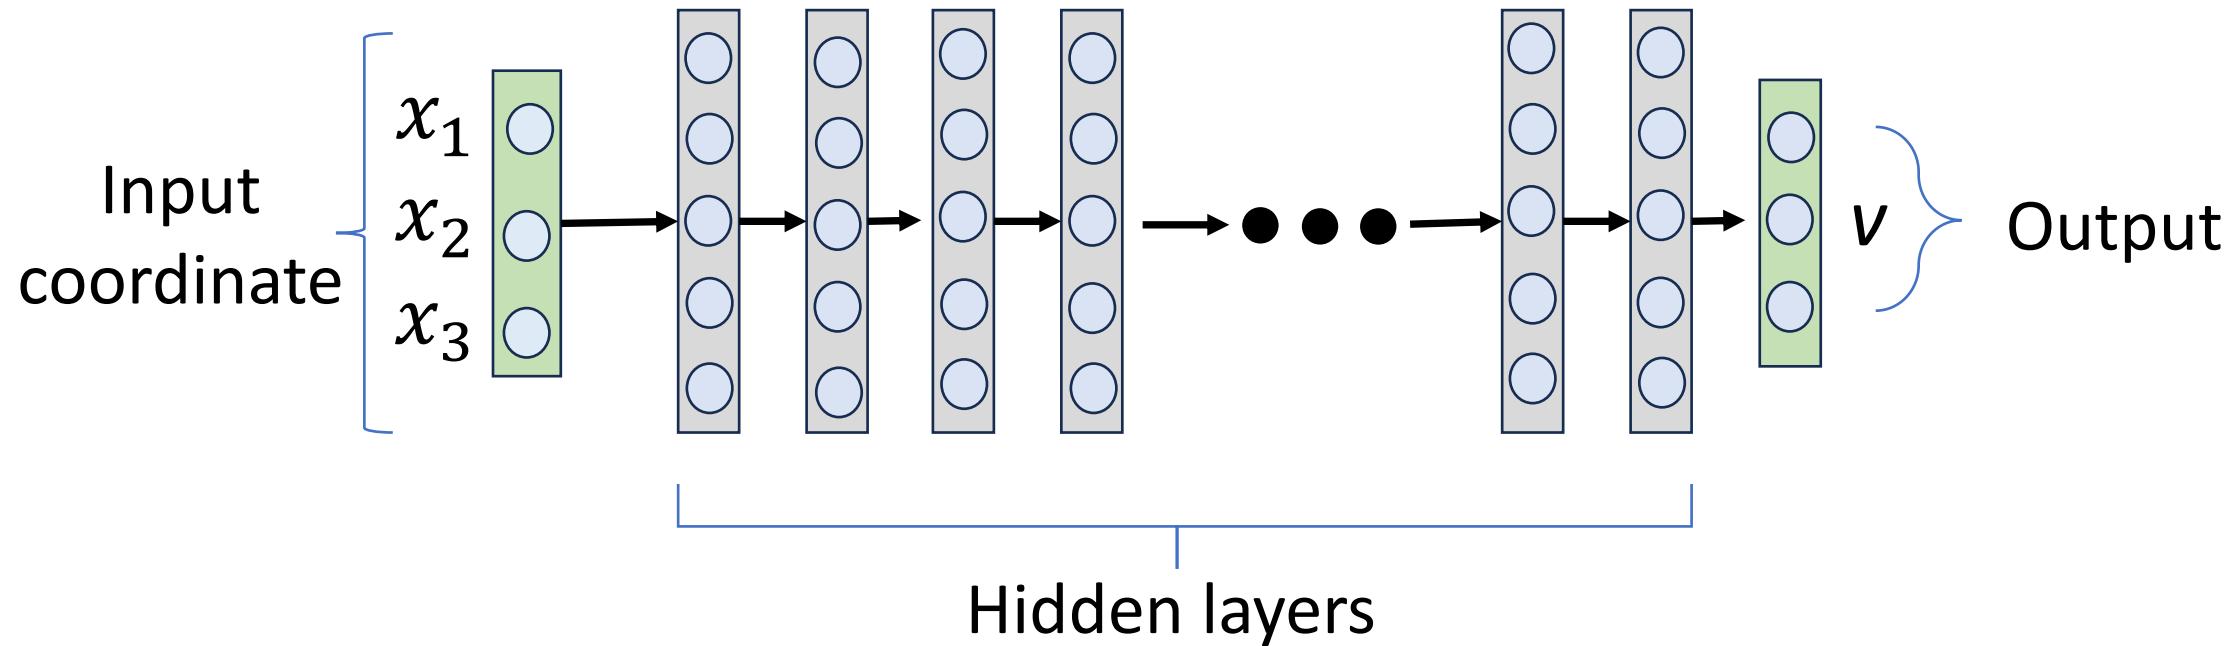


# SIREN: Sinusoidal Representation Network

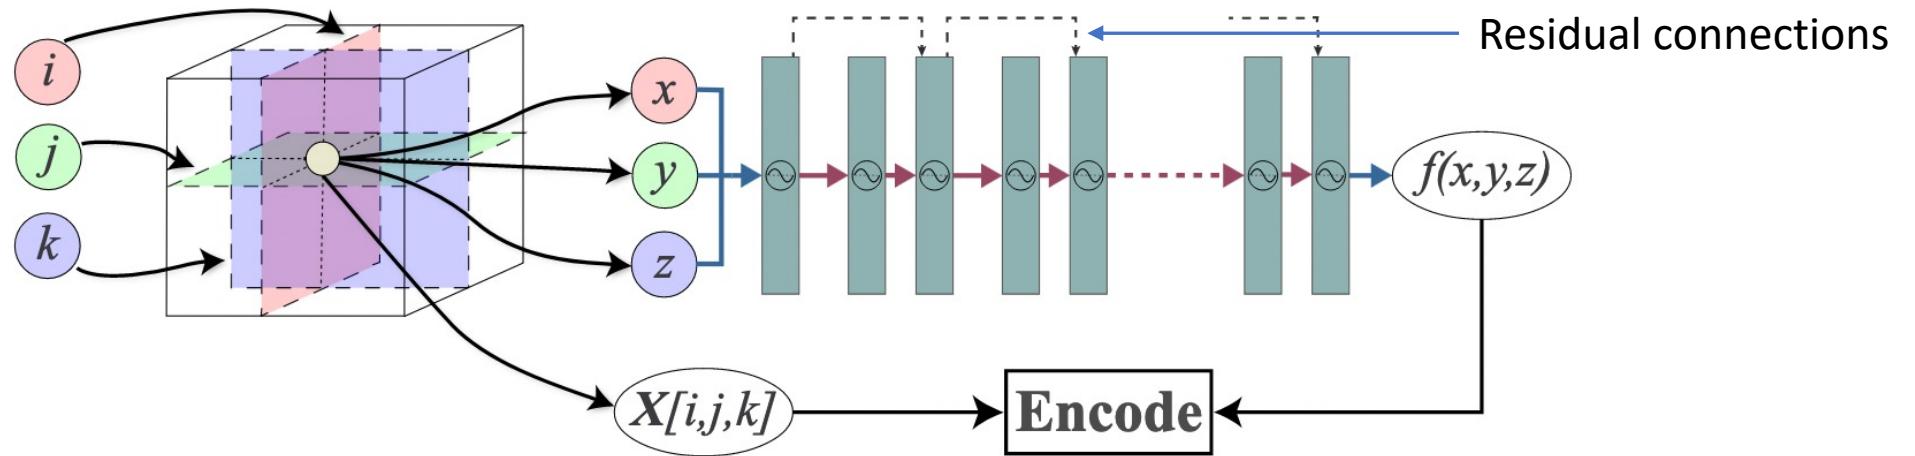


# Compressive Neural Network

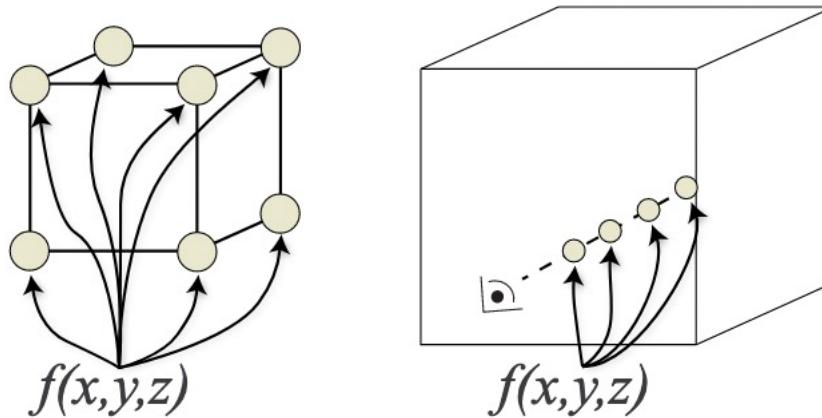
- A Multiplayer Perceptron with sine activation function



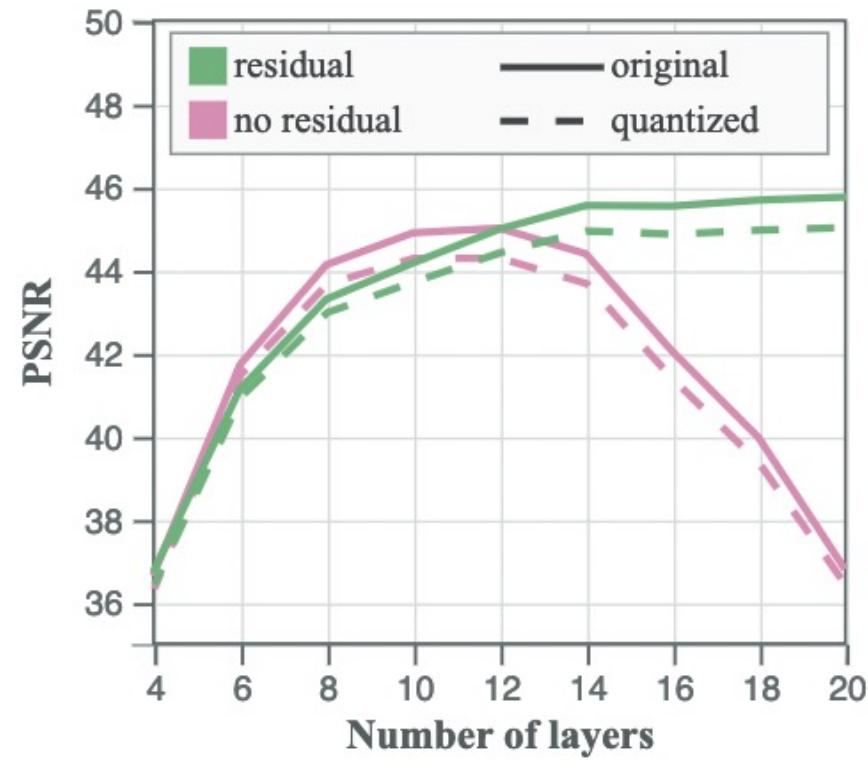
# Compressive Neural Network



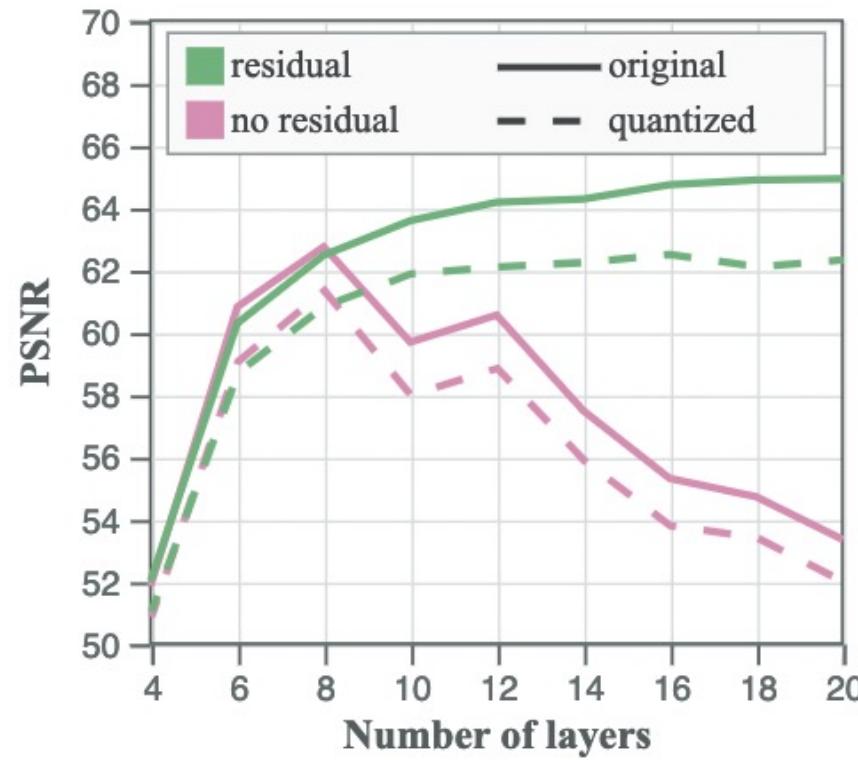
**Decode**



# Benefits of Adding Residual Blocks



(a) Asteroid



(b) Ionization

# Compressive Neural Network

- Loss function Mean squared error (MSE)

$$\min_{\Theta} \sum_{\mathbf{i}} (f_{\Theta}(\mathbf{p_i}) - \mathbf{X}[i_1, i_2, \dots, i_d])^2$$

# Compressive Neural Network

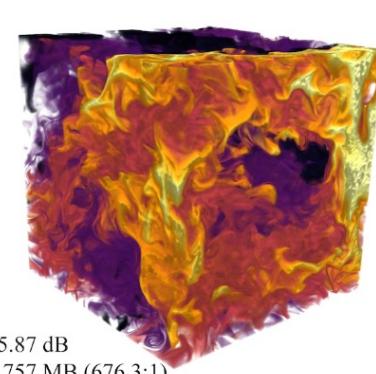
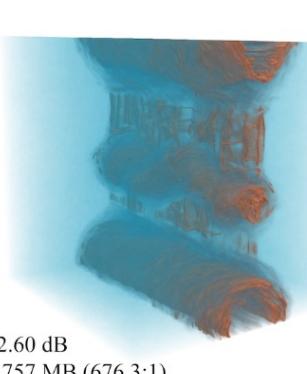
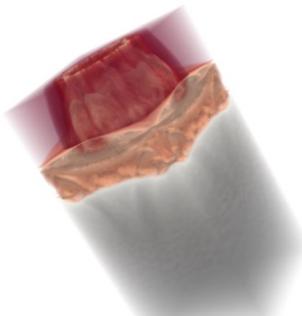
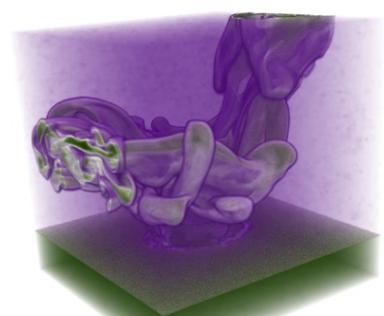
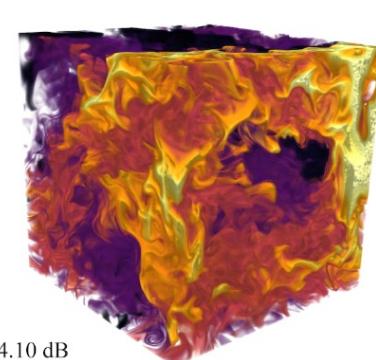
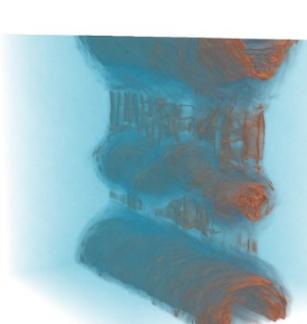
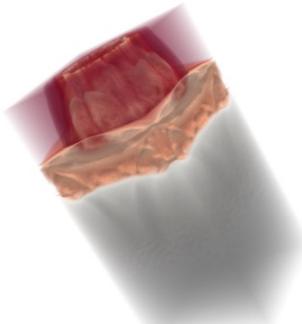
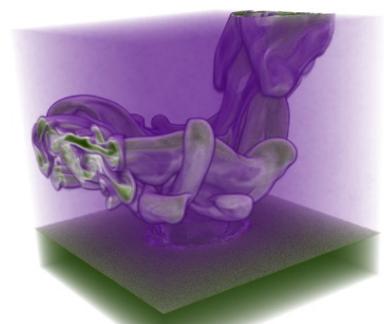
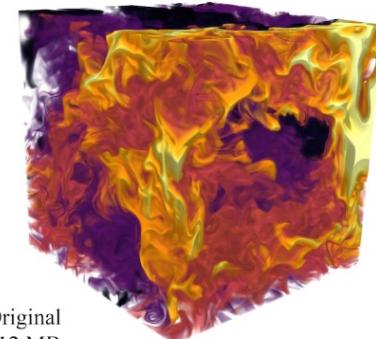
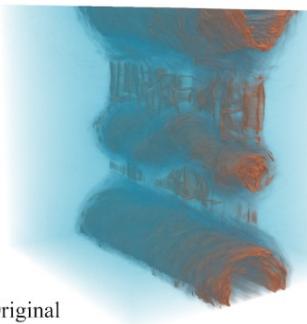
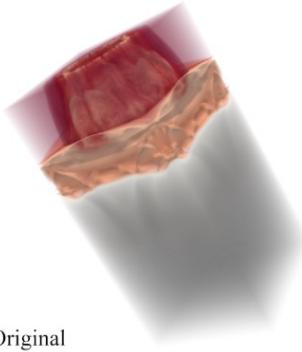
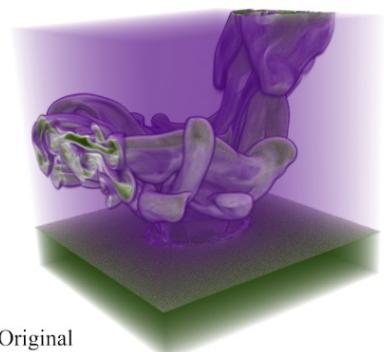
- Loss function Mean squared error (MSE) + Gradient regularization

$$\min_{\Theta} \sum_{\mathbf{i}} (f_{\Theta}(\mathbf{p_i}) - \mathbf{X}[i_1, i_2, \dots, i_d])^2$$

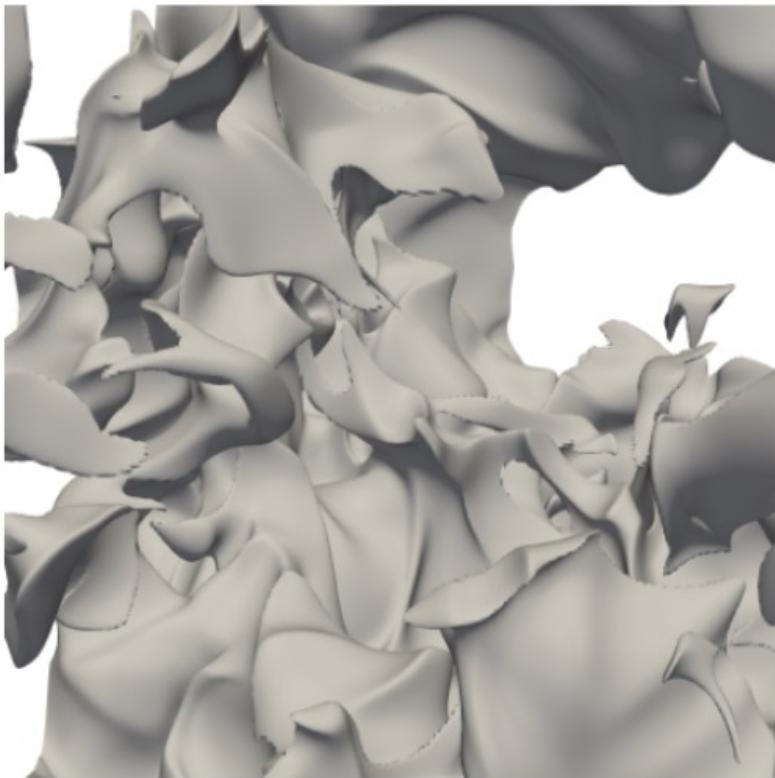


$$\lambda \|\nabla f_{\Theta}(\mathbf{p_i}) - \mathbf{X}'[i_1, i_2, \dots, i_d]\|_2^2$$

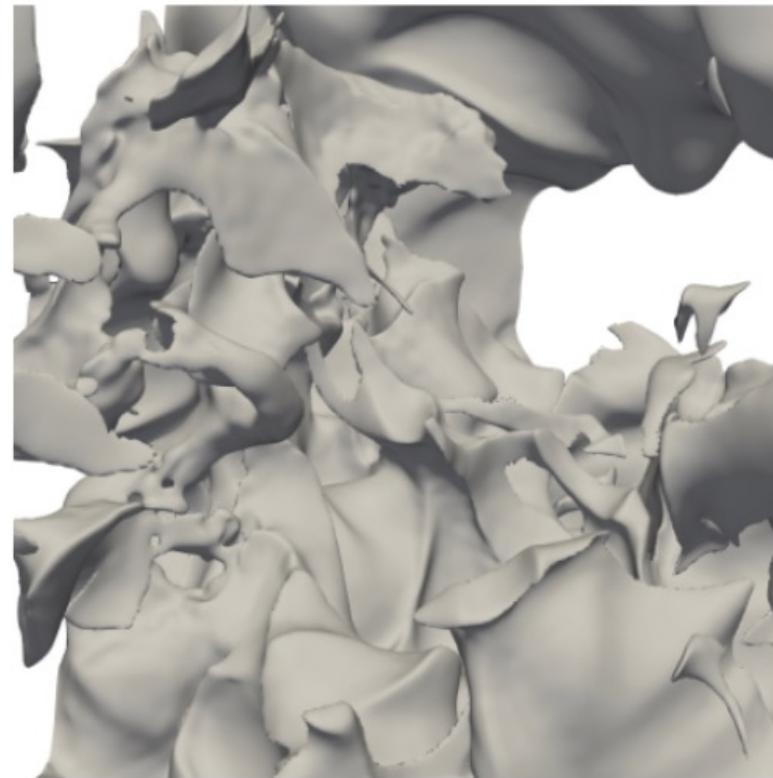
# Reconstruction Results



# Reconstruction Results: Isosurface



(a) ground truth

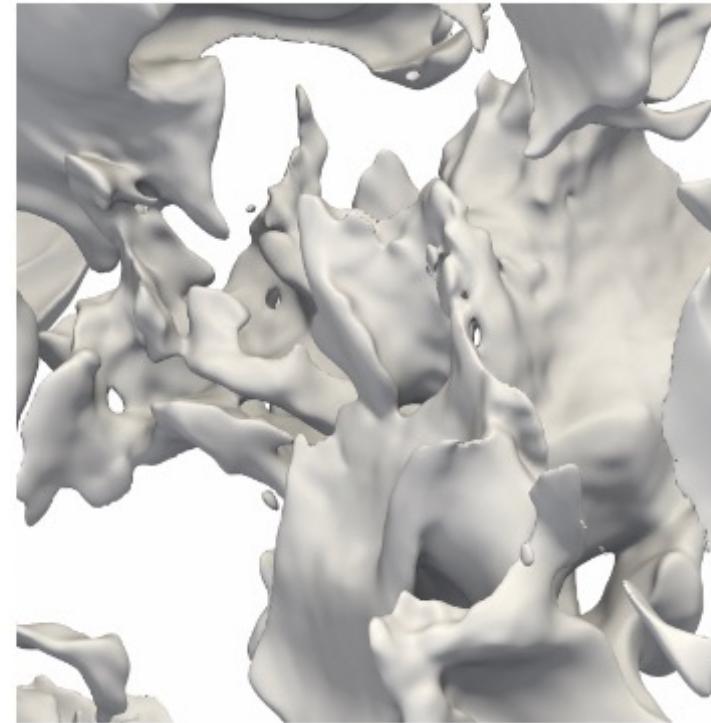


(b) neurcomp

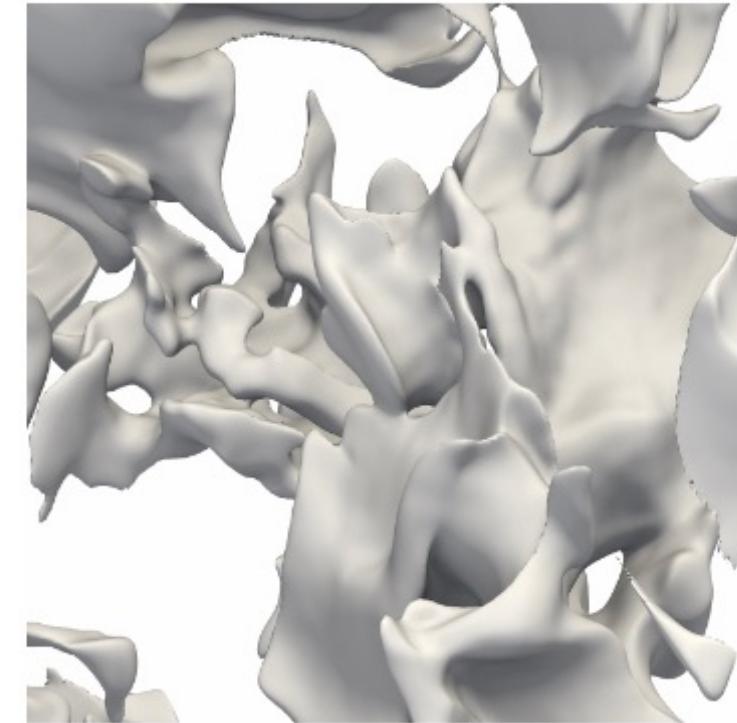
# Impact of Gradient Regularization



(a) ground truth



(b) no gradient reg.



(c) gradient reg.